

# Linear Regression

Shusen Wang

# **Warm-up: Vector and Matrix**

# Vector and Matrix

Vector ( $n$ -dim)

$$\mathbf{a} = \begin{bmatrix} a_1 \\ a_2 \\ \vdots \\ a_n \end{bmatrix}$$

Matrix ( $n \times d$ )

$$\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1d} \\ a_{21} & a_{22} & \cdots & a_{2d} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \cdots & a_{nd} \end{bmatrix}$$


Row and columns

$$\mathbf{A} = \begin{bmatrix} \mathbf{a}_{:1} & \mathbf{a}_{:2} & \cdots & \mathbf{a}_{:d} \end{bmatrix} = \begin{bmatrix} \mathbf{a}_{1:} \\ \mathbf{a}_{2:} \\ \vdots \\ \mathbf{a}_{n:} \end{bmatrix}$$

# Vector Norms

- The  $\ell_p$  norm:  $\|\mathbf{x}\|_p := \left( \sum_i |x_i|^p \right)^{1/p}$ .
- The  $\ell_2$  norm:  $\|\mathbf{x}\|_2 = \left( \sum_i x_i^2 \right)^{1/2}$  (the Euclidean norm).
- The  $\ell_1$  norm  $\|\mathbf{x}\|_1 = \sum_i |x_i|$ .
- The  $\ell_\infty$  norm is defined by  $\|\mathbf{x}\|_\infty = \max_i |x_i|$ .

# Vector Norms



- The  $\ell_2$ -distance (Euclidean distance):  
 $||\mathbf{a} - \mathbf{b}||_2$  (green line)
- The  $\ell_1$ -distance (Manhattan distance):  
 $||\mathbf{a} - \mathbf{b}||_1$  (red, blue, yellow lines)

# Transpose and Rank

Transpose:

$$\begin{bmatrix} 6 & 4 & 24 \\ 1 & -9 & 8 \end{bmatrix}^T = \begin{bmatrix} 6 & 1 \\ 4 & -9 \\ 24 & 8 \end{bmatrix}$$

**Square matrix:** a matrix with the same number of rows and columns.

**Symmetric:** a square matrix  $\mathbf{A}$  is symmetric if  $\mathbf{A}^T = \mathbf{A}$ .

**Rank:** the number of linearly independent rows (or columns).

**Full rank:** a square matrix is full rank if the rank equals to #columns.

# Eigenvalue Decomposition

- Let  $\mathbf{A}$  be any  $n \times n$  symmetric matrix.
- Eigenvalue decomposition:  $\mathbf{A} = \sum_{i=1}^n \lambda_i \mathbf{v}_i \mathbf{v}_i^T$ .
- Eigenvalues satisfy  $|\lambda_1| \geq |\lambda_2| \geq \dots \geq |\lambda_n|$ .
- Eigenvectors satisfy  $\mathbf{v}_i^T \mathbf{v}_j = 0$  for all  $i \neq j$ .
- $\mathbf{A}$  is full rank  $\iff$  all the eigenvalues are nonzero.

# **Warm-up: Optimization**



# Optimization: Basics

Optimization problem:  $\min_{\mathbf{w}} f(\mathbf{w}); \quad \text{s. t. } \mathbf{w} \in \mathcal{C}.$

- $\mathbf{w} = [w_1, \dots, w_d]$  : optimization variables
- $f : \mathbb{R}^d \mapsto \mathbb{R}$  : objective function
- $\mathcal{C}$  (a subset of  $\mathbb{R}^d$ ) : feasible set

# Optimization: Basics

Optimization problem:  $\min_{\mathbf{w}} f(\mathbf{w}) ;$  s. t.  $\mathbf{w} \in \mathcal{C} .$

- $\mathbf{w} = [w_1, \dots, w_d]$  : optimization variables
- $f : \mathbb{R}^d \mapsto \mathbb{R}$  : objective function
- $\mathcal{C}$  (a subset of  $\mathbb{R}^d$ ) : feasible set

  
Constraint

# Optimization: Basics

Optimization problem:  $\min_{\mathbf{w}} f(\mathbf{w}); \quad \text{s. t. } \mathbf{w} \in \mathcal{C}.$

- $\mathbf{w} = [w_1, \dots, w_d]$  : optimization variables
- $f : \mathbb{R}^d \mapsto \mathbb{R}$  : objective function
- $\mathcal{C}$  (a subset of  $\mathbb{R}^d$ ) : feasible set
- $\mathbf{w}^* = \operatorname{argmin}_{\mathbf{w} \in \mathcal{C}} f(\mathbf{w})$  is the optimal solution to the problem
  - $f(\mathbf{w}^*) \leq f(\mathbf{w})$  for all the vectors  $\mathbf{w}$  in the set  $\mathcal{C}$ .
  - $\mathbf{w}^*$  may not exist; if it exists, it may not be unique.

# Least Squares Regression

# The Linear Regression Task


**Input:** vectors  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  and labels  $y_1, \dots, y_n \in \mathbb{R}$

**Output:** a vector  $\mathbf{w} \in \mathbb{R}^d$  and scalar  $b \in \mathbb{R}$  such that  $\mathbf{x}_i^T \mathbf{w} + b \approx y_i$ .

1-dim ( $d = 1$ ) example:

Solution:

$$y_i \approx 0.15 x_i + 5.0$$




# The Linear Regression Task

**Input:** vectors  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  and labels  $y_1, \dots, y_n \in \mathbb{R}$

**Output:** a vector  $\mathbf{w} \in \mathbb{R}^d$  and scalar  $b \in \mathbb{R}$  such that  $\mathbf{x}_i^T \mathbf{w} + b \approx y_i$ .

**Question** (regard training):  
how to compute  $\mathbf{w}$  and  $b$ ?



# The Linear Regression Task

**Input:** vectors  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  and labels  $y_1, \dots, y_n \in \mathbb{R}$

**Output:** a vector  $\mathbf{w} \in \mathbb{R}^d$  and scalar  $b \in \mathbb{R}$  such that  $\mathbf{x}_i^T \mathbf{w} + b \approx y_i$ .

**Method:** least squares regression.

- The optimization model:

$$\min_{\mathbf{w}, b} L(\mathbf{w}, b), \quad \text{where } L(\mathbf{w}, b) = \sum_{i=1}^n (\mathbf{x}_i^T \mathbf{w} + b - y_i)^2$$

# Least Squares Regression

- The optimization model:


$$\min_{\mathbf{w}, b} L(\mathbf{w}, b), \quad \text{where } L(\mathbf{w}, b) = \sum_{i=1}^n (\mathbf{x}_i^T \mathbf{w} + b - y_i)^2$$



# Least Squares Regression

- The optimization model:

$$\min_{\mathbf{w}, b} L(\mathbf{w}, b), \quad \text{where } L(\mathbf{w}, b) = \sum_{i=1}^n (\mathbf{x}_i^T \mathbf{w} + \boxed{b} - y_i)^2$$

  
Intercept

# Least Squares Regression

- The optimization model:

$$\min_{\mathbf{w}, b} L(\mathbf{w}, b), \quad \text{where } L(\mathbf{w}, b) = \sum_{i=1}^n (\mathbf{x}_i^T \mathbf{w} + b - y_i)^2$$



$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \sum_{i=1}^n (\bar{\mathbf{x}}_i^T \bar{\mathbf{w}} - y_i)^2$$

- Define  $\bar{\mathbf{x}}_i = [\mathbf{x}_i; 1] \in \mathbb{R}^{d+1}$
- Define  $\bar{\mathbf{w}} = [\mathbf{w}, b] \in \mathbb{R}^{d+1}$
- $\Rightarrow \mathbf{x}_i^T \mathbf{w} + b = \bar{\mathbf{x}}_i^T \bar{\mathbf{w}}$

# Least Squares Regression

- The optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \sum_{i=1}^n (\bar{\mathbf{x}}_i^T \bar{\mathbf{w}} - y_i)^2$$



Matrix form:  $\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \|\bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y}\|_2^2$

$$\bar{\mathbf{X}} = \begin{bmatrix} \mathbf{x}_1^T & 1 \\ \mathbf{x}_2^T & 1 \\ \vdots & \vdots \\ \mathbf{x}_n^T & 1 \end{bmatrix}$$

$$n \times (d + 1)$$

$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

$$n \times 1$$

$$\bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} = \begin{bmatrix} \bar{\mathbf{x}}_1^T \bar{\mathbf{w}} - y_1 \\ \bar{\mathbf{x}}_2^T \bar{\mathbf{w}} - y_2 \\ \vdots \\ \bar{\mathbf{x}}_n^T \bar{\mathbf{w}} - y_n \end{bmatrix}$$

$$n \times 1$$

# Least Squares Regression

- The optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2$$

Tasks

Linear  
Regression

Methods

Least Squares Regression

LASSO

Least Absolute Deviations

Algorithms

?

# Least Squares Regression

- The optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2$$

Tasks

Linear  
Regression

Methods

Least Squares Regression

LASSO

Least Absolute Deviations

Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient (CG)

# Least Squares Regression

- Solve the optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2$$

Gradient:  $\frac{\partial \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2}{\partial \bar{\mathbf{w}}} = 2(\bar{\mathbf{X}}^T \bar{\mathbf{X}} \bar{\mathbf{w}} - \bar{\mathbf{X}}^T \mathbf{y})$

## Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient

# Least Squares Regression

- Solve the optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2$$

Gradient:  $\frac{\partial \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2}{\partial \bar{\mathbf{w}}} = 2(\bar{\mathbf{X}}^T \bar{\mathbf{X}} \bar{\mathbf{w}} - \bar{\mathbf{X}}^T \mathbf{y}) = \mathbf{0}$

1<sup>st</sup>-order optimal condition

**Algorithms**

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient

# Least Squares Regression

- Solve the optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2$$

Gradient:  $\frac{\partial \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2}{\partial \bar{\mathbf{w}}} = 2(\bar{\mathbf{X}}^T \bar{\mathbf{X}} \bar{\mathbf{w}} - \bar{\mathbf{X}}^T \mathbf{y}) = \mathbf{0}$



Normal equation:  $\bar{\mathbf{X}}^T \bar{\mathbf{X}} \bar{\mathbf{w}} = \bar{\mathbf{X}}^T \mathbf{y}$

Assume  $\bar{\mathbf{X}}^T \bar{\mathbf{X}}$  is full rank.



Analytical solution:  $\bar{\mathbf{w}}^* = (\bar{\mathbf{X}}^T \bar{\mathbf{X}})^{-1} \bar{\mathbf{X}}^T \mathbf{y}$

## Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient



# Least Squares Regression

- Solve the optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2$$

Gradient:  $\frac{\partial \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2}{\partial \bar{\mathbf{w}}} = 2(\bar{\mathbf{X}}^T \bar{\mathbf{X}} \bar{\mathbf{w}} - \bar{\mathbf{X}}^T \mathbf{y}) = \mathbf{0}$

Gradient descent repeats:

1. Compute gradient:  $\mathbf{g}_t = \bar{\mathbf{X}}^T \bar{\mathbf{X}} \bar{\mathbf{w}}_t - \bar{\mathbf{X}}^T \mathbf{y}$
2. Update:  $\bar{\mathbf{w}}_{t+1} = \bar{\mathbf{w}}_t - \alpha_t \mathbf{g}_t$

**Algorithms**

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient

# Least Squares Regression

- Solve the optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2$$

Convergence: after  $O\left(\kappa \log \frac{1}{\epsilon}\right)$  iterations,

$$\left\| \bar{\mathbf{X}} (\bar{\mathbf{w}}_t - \bar{\mathbf{w}}^*) \right\|_2 \leq \epsilon \left\| \bar{\mathbf{X}} (\bar{\mathbf{w}}_0 - \bar{\mathbf{w}}^*) \right\|_2.$$

$$\kappa = \frac{\lambda_{\max}(\bar{\mathbf{X}}^T \bar{\mathbf{X}})}{\lambda_{\min}(\bar{\mathbf{X}}^T \bar{\mathbf{X}})} \text{ is the condition number.}$$

## Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient

# Least Squares Regression

- Solve the optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2$$

Convergence: after  $O\left(\sqrt{\kappa} \log \frac{1}{\epsilon}\right)$  iterations,

$$\left\| \bar{\mathbf{X}} (\bar{\mathbf{w}}_t - \bar{\mathbf{w}}^*) \right\|_2 \leq \epsilon \left\| \bar{\mathbf{X}} (\bar{\mathbf{w}}_0 - \bar{\mathbf{w}}^*) \right\|_2.$$

$$\kappa = \frac{\lambda_{\max}(\bar{\mathbf{X}}^T \bar{\mathbf{X}})}{\lambda_{\min}(\bar{\mathbf{X}}^T \bar{\mathbf{X}})} \text{ is the condition number.}$$

The pseudo-code of CG is available at the [Wikipedia](#).

## Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient

# Least Squares Regression

- Solve the optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2$$

Tasks

Linear  
Regression

Methods

Least Squares Regression

LASSO

Least Absolute Deviations

Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient

# **Solve Least Squares in Python**

# 1. Load Data

```
from keras.datasets import boston_housing

(x_train, y_train), (x_test, y_test) = boston_housing.load_data()

print('shape of x_train: ' + str(x_train.shape))
print('shape of x_test: ' + str(x_test.shape))
print('shape of y_train: ' + str(y_train.shape))
print('shape of y_test: ' + str(y_test.shape))
```

```
shape of x_train: (404, 13)
shape of x_test: (102, 13)
shape of y_train: (404,)
shape of y_test: (102,)
```

## 2. Add A Feature

```
import numpy
```

```
n, d = x_train.shape
```

```
xbar_train = numpy.concatenate((x_train, numpy.ones((n, 1))),  
                                axis=1)
```

```
print('shape of x_train: ' + str(x_train.shape))
```

```
print('shape of xbar_train: ' + str(xbar_train.shape))
```

```
shape of x_train: (404, 13)
```

```
shape of xbar_train: (404, 14)
```

### 3. Solve the Least Squares

Analytical solution:  $\bar{\mathbf{w}} = (\bar{\mathbf{X}}^T \bar{\mathbf{X}})^{-1} \bar{\mathbf{X}}^T \mathbf{y}$

```
xx = numpy.dot(xbar_train.T, xbar_train)
xx_inv = numpy.linalg.pinv(xx)
xy = numpy.dot(xbar_train.T, y_train)
w = numpy.dot(xx_inv, xy)
```



### 3. Solve the Least Squares

Analytical solution:  $\bar{\mathbf{w}} = (\bar{\mathbf{X}}^T \bar{\mathbf{X}})^{-1} \bar{\mathbf{X}}^T \mathbf{y}$

```
xx = numpy.dot(xbar_train.T, xbar_train)
xx_inv = numpy.linalg.pinv(xx)
xy = numpy.dot(xbar_train.T, y_train)
w = numpy.dot(xx_inv, xy)
```

### 3. Solve the Least Squares

Analytical solution:  $\bar{\mathbf{w}} = (\bar{\mathbf{X}}^T \bar{\mathbf{X}})^{-1} \bar{\mathbf{X}}^T \mathbf{y}$

```
xx = numpy.dot(xbar_train.T, xbar_train)
xx_inv = numpy.linalg.pinv(xx)
xy = numpy.dot(xbar_train.T, y_train)
w = numpy.dot(xx_inv, xy)
```

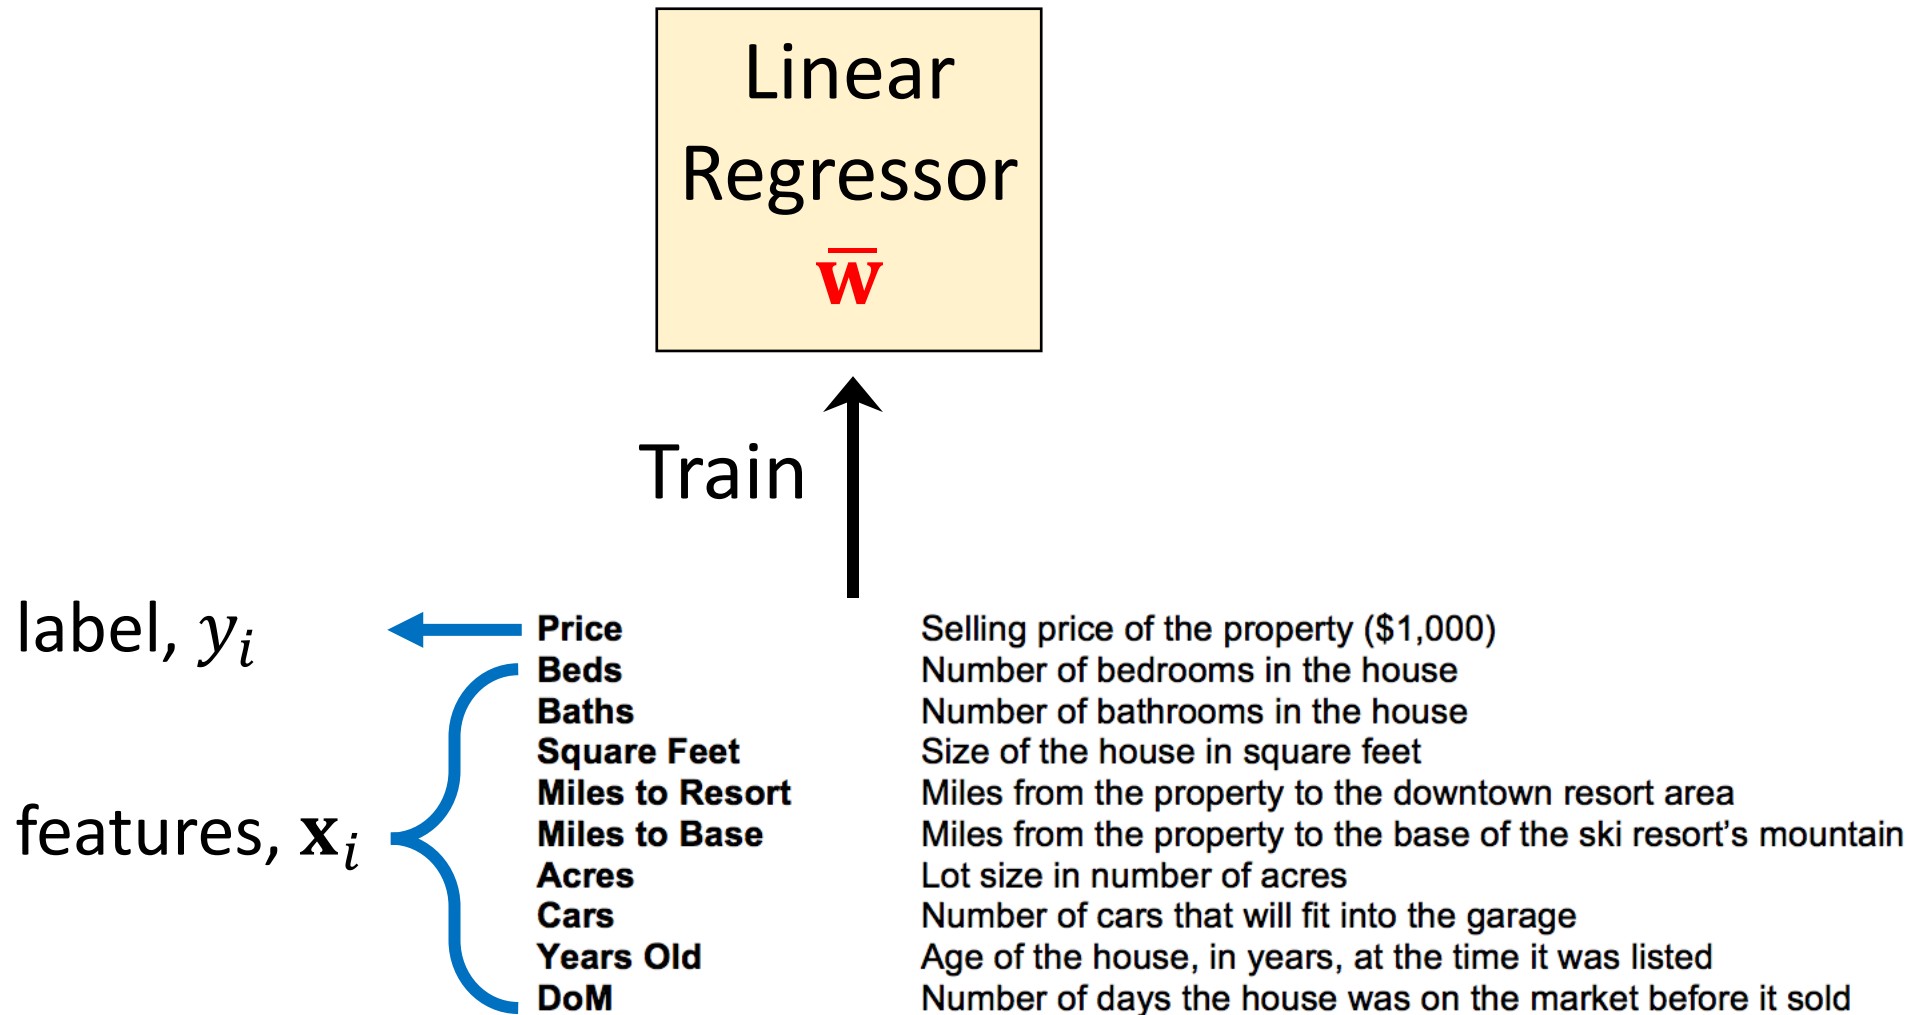
### 3. Solve the Least Squares

Training Mean Squared Error (MSE):  $\frac{1}{n} \|\mathbf{y} - \bar{\mathbf{X}}\bar{\mathbf{w}}\|_2^2$

```
y_lsr = numpy.dot(xbar_train, w)
diff = y_lsr - y_train
mse = numpy.mean(diff * diff)
print('Train MSE: ' + str(mse))
```

Train MSE: 22.00480083834814

# Linear Regression for Housing Price



# Linear Regression for Housing Price



Features of a House,  $\mathbf{x}'$   
→ Extend it to  $\bar{\mathbf{x}}'$

Linear  
Regressor  
 $\bar{\mathbf{w}}$

Predict

Price:  
 $\bar{\mathbf{w}}^T \bar{\mathbf{x}}' = \$500\text{K}$

Train

label,  $y_i$

features,  $\mathbf{x}_i$

Price

Beds

Baths

Square Feet

Miles to Resort

Miles to Base

Acres

Cars

Years Old

DoM

Selling price of the property (\$1,000)

Number of bedrooms in the house

Number of bathrooms in the house

Size of the house in square feet

Miles from the property to the downtown resort area

Miles from the property to the base of the ski resort's mountain

Lot size in number of acres

Number of cars that will fit into the garage

Age of the house, in years, at the time it was listed

Number of days the house was on the market before it sold

## 4. Make Prediction for Test Samples

- Add a feature to the test feature matrix:  $\mathbf{X}_{\text{test}} \rightarrow \bar{\mathbf{X}}_{\text{test}}$ .
- Make prediction by:  $\mathbf{y}_{\text{pred}} = \bar{\mathbf{X}}_{\text{test}} \bar{\mathbf{w}}$ .

```
n_test, _ = x_test.shape
xbar_test = numpy.concatenate((x_test, numpy.ones((n_test, 1))), axis=1)
y_pred = numpy.dot(xbar_test, w)
```

# 4. Make Prediction for Test Samples

- Add a feature to the test feature matrix:  $\mathbf{X}_{\text{test}} \rightarrow \bar{\mathbf{X}}_{\text{test}}$ .
- Make prediction by:  $\mathbf{y}_{\text{pred}} = \bar{\mathbf{X}}_{\text{test}} \bar{\mathbf{w}}$ .
- MSE (testing):  $\frac{1}{n_{\text{test}}} \left\| \mathbf{y}_{\text{pred}} - \mathbf{y}_{\text{test}} \right\|_2^2$

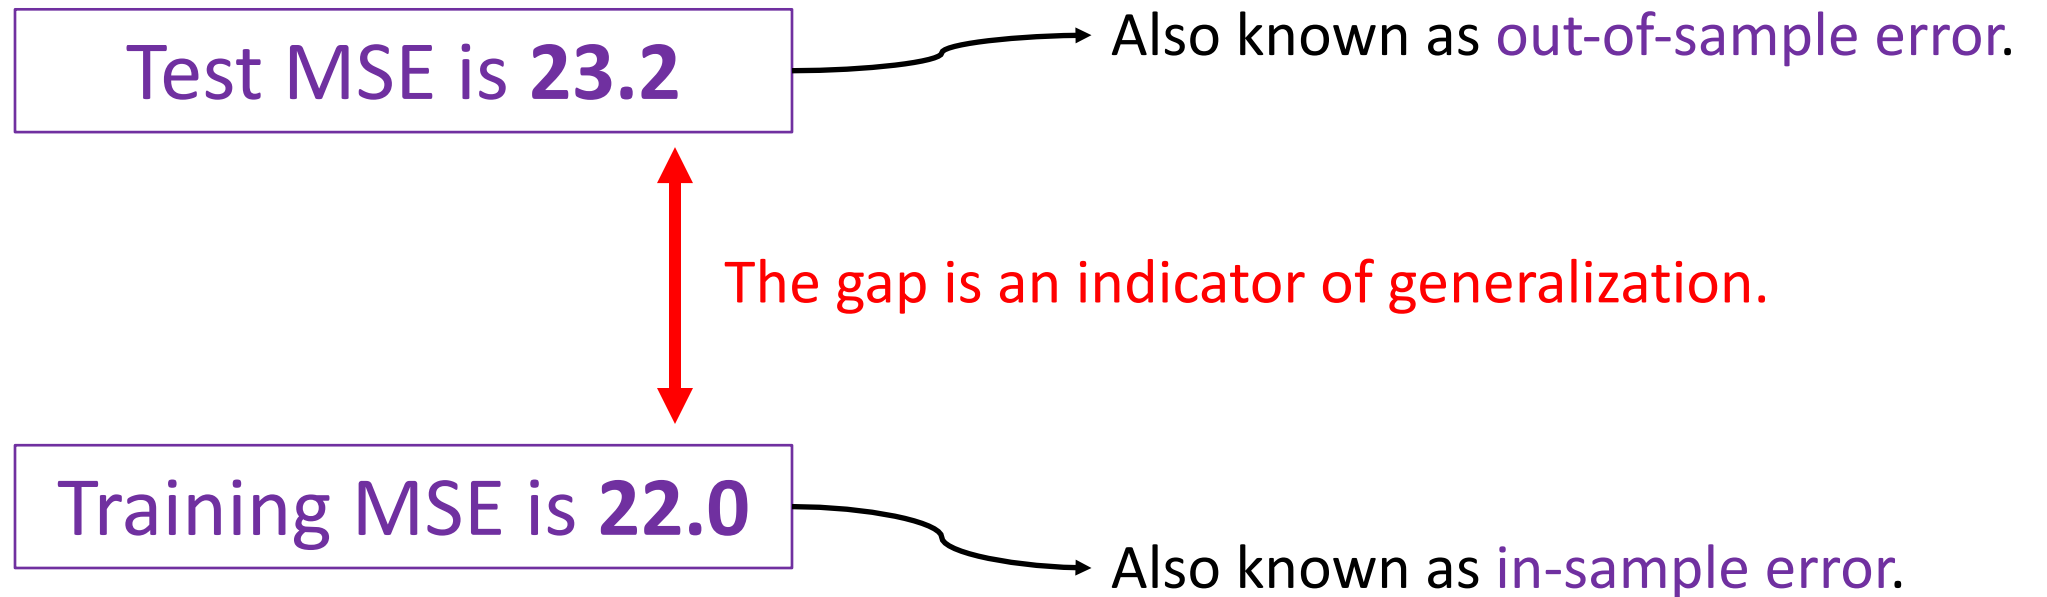
```
# mean squared error (testing)
```

```
diff = y_pred - y_test  
mse = numpy.mean(diff * diff)  
print('Test MSE: ' + str(mse))
```

Test MSE: 23.195599256409857

Training MSE is **22.0**

# 4. Make Prediction for Testing Samples





# 5. Compare with Baseline

Baseline:

- whatever the features are, the prediction is  $\text{mean}(\mathbf{y})$ .

```
y_mean = numpy.mean(y_train)

diff = y_pred - y_mean
mse = numpy.mean(diff * diff)
print('Test MSE: ' + str(mse))
```

Test MSE: 57.38297638530044

Test MSE of least  
squares is **23.19**