

Linear Regression

Shusen Wang

Warm-up: Vector and Matrix

Vector and Matrix

Vector (n -dim)

$$\mathbf{a} = \begin{bmatrix} a_1 \\ a_2 \\ \vdots \\ a_n \end{bmatrix}$$

Matrix ($n \times d$)

$$\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1d} \\ a_{21} & a_{22} & \cdots & a_{2d} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \cdots & a_{nd} \end{bmatrix}$$


Row and columns

$$\mathbf{A} = \begin{bmatrix} \mathbf{a}_{:1} & \mathbf{a}_{:2} & \cdots & \mathbf{a}_{:d} \end{bmatrix} = \begin{bmatrix} \mathbf{a}_{1:} \\ \mathbf{a}_{2:} \\ \vdots \\ \mathbf{a}_{n:} \end{bmatrix}$$

Vector Norms

- The ℓ_p norm: $\|\mathbf{x}\|_p := \left(\sum_i |x_i|^p \right)^{1/p}$.
- The ℓ_2 norm: $\|\mathbf{x}\|_2 = \left(\sum_i x_i^2 \right)^{1/2}$ (the Euclidean norm).
- The ℓ_1 norm $\|\mathbf{x}\|_1 = \sum_i |x_i|$.
- The ℓ_∞ norm is defined by $\|\mathbf{x}\|_\infty = \max_i |x_i|$.

Vector Norms



- The ℓ_2 -distance (Euclidean distance):
 $||\mathbf{a} - \mathbf{b}||_2$ (green line)
- The ℓ_1 -distance (Manhattan distance):
 $||\mathbf{a} - \mathbf{b}||_1$ (red, blue, yellow lines)

Transpose and Rank

Transpose:

$$\begin{bmatrix} 6 & 4 & 24 \\ 1 & -9 & 8 \end{bmatrix}^T = \begin{bmatrix} 6 & 1 \\ 4 & -9 \\ 24 & 8 \end{bmatrix}$$

Square matrix: a matrix with the same number of rows and columns.

Symmetric: a square matrix \mathbf{A} is symmetric if $\mathbf{A}^T = \mathbf{A}$.

Rank: the number of linearly independent rows (or columns).

Full rank: a square matrix is full rank if the rank equals to #columns.

Eigenvalue Decomposition

- Let \mathbf{A} be any $n \times n$ symmetric matrix.
- Eigenvalue decomposition: $\mathbf{A} = \sum_{i=1}^n \lambda_i \mathbf{v}_i \mathbf{v}_i^T$.
- Eigenvalues satisfy $|\lambda_1| \geq |\lambda_2| \geq \dots \geq |\lambda_n|$.
- Eigenvectors satisfy $\mathbf{v}_i^T \mathbf{v}_j = 0$ for all $i \neq j$.
- \mathbf{A} is full rank \iff all the eigenvalues are nonzero.

Warm-up: Optimization

Optimization: Basics

Optimization problem: $\min_{\mathbf{w}} f(\mathbf{w}); \quad \text{s. t. } \mathbf{w} \in \mathcal{C}.$

- $\mathbf{w} = [w_1, \dots, w_d]$: optimization variables
- $f : \mathbb{R}^d \mapsto \mathbb{R}$: objective function
- \mathcal{C} (a subset of \mathbb{R}^d) : feasible set

Optimization: Basics

Optimization problem: $\min_{\mathbf{w}} f(\mathbf{w}) ;$ s. t. $\mathbf{w} \in \mathcal{C}$.

- $\mathbf{w} = [w_1, \dots, w_d]$: optimization variables
- $f : \mathbb{R}^d \mapsto \mathbb{R}$: objective function
- \mathcal{C} (a subset of \mathbb{R}^d) : feasible set


Constraint

Optimization: Basics

Optimization problem: $\min_{\mathbf{w}} f(\mathbf{w}); \quad \text{s. t. } \mathbf{w} \in \mathcal{C}.$

- $\mathbf{w}^* = \operatorname{argmin}_{\mathbf{w} \in \mathcal{C}} f(\mathbf{w})$ is the optimal solution to the problem
- $f(\mathbf{w}^*) \leq f(\mathbf{w})$ for all the vectors \mathbf{w} in the set \mathcal{C} .
- \mathbf{w}^* may not exist, e.g., \mathcal{C} is the empty set.
- If \mathbf{w}^* exists, it may not be unique.

Least Squares Regression

The Linear Regression Task


Input: vectors $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$ and labels $y_1, \dots, y_n \in \mathbb{R}$

Output: a vector $\mathbf{w} \in \mathbb{R}^d$ and scalar $b \in \mathbb{R}$ such that $\mathbf{x}_i^T \mathbf{w} + b \approx y_i$.

1-dim ($d = 1$) example:

Solution:

$$y_i \approx 0.15 x_i + 5.0$$




The Linear Regression Task

Input: vectors $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$ and labels $y_1, \dots, y_n \in \mathbb{R}$

Output: a vector $\mathbf{w} \in \mathbb{R}^d$ and scalar $b \in \mathbb{R}$ such that $\mathbf{x}_i^T \mathbf{w} + b \approx y_i$.

Question (regard training):
how to compute \mathbf{w} and b ?



The Linear Regression Task

Input: vectors $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$ and labels $y_1, \dots, y_n \in \mathbb{R}$

Output: a vector $\mathbf{w} \in \mathbb{R}^d$ and scalar $b \in \mathbb{R}$ such that $\mathbf{x}_i^T \mathbf{w} + b \approx y_i$.

Method: least squares regression.

- The optimization model:

$$\min_{\mathbf{w}, b} L(\mathbf{w}, b), \quad \text{where } L(\mathbf{w}, b) = \sum_{i=1}^n (\mathbf{x}_i^T \mathbf{w} + b - y_i)^2$$

Least Squares Regression


- The optimization model:

$$\min_{\mathbf{w}, b} L(\mathbf{w}, b), \quad \text{where } L(\mathbf{w}, b) = \sum_{i=1}^n (\mathbf{x}_i^T \mathbf{w} + b - y_i)^2$$

Least Squares Regression

- The optimization model:

$$\min_{\mathbf{w}, b} L(\mathbf{w}, b), \quad \text{where } L(\mathbf{w}, b) = \sum_{i=1}^n (\mathbf{x}_i^T \mathbf{w} + \boxed{b} - y_i)^2$$


Intercept

Least Squares Regression

- The optimization model:

$$\min_{\mathbf{w}, b} L(\mathbf{w}, b), \quad \text{where } L(\mathbf{w}, b) = \sum_{i=1}^n (\mathbf{x}_i^T \mathbf{w} + b - y_i)^2$$



$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \sum_{i=1}^n (\bar{\mathbf{x}}_i^T \bar{\mathbf{w}} - y_i)^2$$

- Define $\bar{\mathbf{x}}_i = [\mathbf{x}_i; 1] \in \mathbb{R}^{d+1}$
- Define $\bar{\mathbf{w}} = [\mathbf{w}, b] \in \mathbb{R}^{d+1}$
- $\Rightarrow \mathbf{x}_i^T \mathbf{w} + b = \bar{\mathbf{x}}_i^T \bar{\mathbf{w}}$

Least Squares Regression

- The optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \sum_{i=1}^n (\bar{\mathbf{x}}_i^T \bar{\mathbf{w}} - y_i)^2$$

$$\bar{\mathbf{X}} = \begin{bmatrix} \mathbf{x}_1^T & 1 \\ \mathbf{x}_2^T & 1 \\ \vdots & \vdots \\ \mathbf{x}_n^T & 1 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

$$n \times (d + 1)$$

$$n \times 1$$

Least Squares Regression

- The optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \sum_{i=1}^n (\bar{\mathbf{x}}_i^T \bar{\mathbf{w}} - y_i)^2$$



Matrix form:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \|\bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y}\|_2^2$$

$$\bar{\mathbf{X}} = \begin{bmatrix} \mathbf{x}_1^T & 1 \\ \mathbf{x}_2^T & 1 \\ \vdots & \vdots \\ \mathbf{x}_n^T & 1 \end{bmatrix}$$

$n \times (d + 1)$

$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

$n \times 1$

$$\bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} = \begin{bmatrix} \bar{\mathbf{x}}_1^T \bar{\mathbf{w}} - y_1 \\ \bar{\mathbf{x}}_2^T \bar{\mathbf{w}} - y_2 \\ \vdots \\ \bar{\mathbf{x}}_n^T \bar{\mathbf{w}} - y_n \end{bmatrix}$$

$n \times 1$

Least Squares Regression

- The optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2$$

Tasks

Linear
Regression

Methods

Least Squares Regression

LASSO

Least Absolute Deviations

Algorithms

?

Least Squares Regression

- The optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2$$

Tasks

Linear
Regression

Methods

Least Squares Regression

LASSO

Least Absolute Deviations

Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient (CG)

Least Squares Regression

- Solve the optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2$$

Gradient: $\frac{\partial \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2}{\partial \bar{\mathbf{w}}} = 2(\bar{\mathbf{X}}^T \bar{\mathbf{X}} \bar{\mathbf{w}} - \bar{\mathbf{X}}^T \mathbf{y})$

Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient

Least Squares Regression

- Solve the optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2$$

Gradient: $\frac{\partial \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2}{\partial \bar{\mathbf{w}}} = 2(\bar{\mathbf{X}}^T \bar{\mathbf{X}} \bar{\mathbf{w}} - \bar{\mathbf{X}}^T \mathbf{y}) = \mathbf{0}$

1st-order optimality condition

Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient

Least Squares Regression

- Solve the optimization model:


$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2$$

Gradient: $\frac{\partial \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2}{\partial \bar{\mathbf{w}}} = 2(\bar{\mathbf{X}}^T \bar{\mathbf{X}} \bar{\mathbf{w}} - \bar{\mathbf{X}}^T \mathbf{y}) = \mathbf{0}$



Normal equation: $\bar{\mathbf{X}}^T \bar{\mathbf{X}} \bar{\mathbf{w}} = \bar{\mathbf{X}}^T \mathbf{y}$

Assume $\bar{\mathbf{X}}^T \bar{\mathbf{X}}$ is full rank.



Analytical solution: $\bar{\mathbf{w}}^* = (\bar{\mathbf{X}}^T \bar{\mathbf{X}})^{-1} \bar{\mathbf{X}}^T \mathbf{y}$

Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient

Least Squares Regression

- Solve the optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2$$

Gradient: $\frac{\partial \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2}{\partial \bar{\mathbf{w}}} = 2(\bar{\mathbf{X}}^T \bar{\mathbf{X}} \bar{\mathbf{w}} - \bar{\mathbf{X}}^T \mathbf{y}) = \mathbf{0}$

Gradient descent repeats:

1. Compute gradient: $\mathbf{g}_t = \bar{\mathbf{X}}^T \bar{\mathbf{X}} \bar{\mathbf{w}}_t - \bar{\mathbf{X}}^T \mathbf{y}$
2. Update: $\bar{\mathbf{w}}_{t+1} = \bar{\mathbf{w}}_t - \alpha_t \mathbf{g}_t$

Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient

Least Squares Regression

- Solve the optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2$$

Convergence: after $O\left(\kappa \log \frac{1}{\epsilon}\right)$ iterations,

$$\left\| \bar{\mathbf{X}} (\bar{\mathbf{w}}_t - \bar{\mathbf{w}}^*) \right\|_2 \leq \epsilon \left\| \bar{\mathbf{X}} (\bar{\mathbf{w}}_0 - \bar{\mathbf{w}}^*) \right\|_2.$$

$$\kappa = \frac{\lambda_{\max}(\bar{\mathbf{X}}^T \bar{\mathbf{X}})}{\lambda_{\min}(\bar{\mathbf{X}}^T \bar{\mathbf{X}})} \text{ is the condition number.}$$

Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient

Least Squares Regression

- Solve the optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2$$

Convergence: after $O\left(\sqrt{\kappa} \log \frac{1}{\epsilon}\right)$ iterations,

$$\left\| \bar{\mathbf{X}} (\bar{\mathbf{w}}_t - \bar{\mathbf{w}}^*) \right\|_2 \leq \epsilon \left\| \bar{\mathbf{X}} (\bar{\mathbf{w}}_0 - \bar{\mathbf{w}}^*) \right\|_2.$$

$$\kappa = \frac{\lambda_{\max}(\bar{\mathbf{X}}^T \bar{\mathbf{X}})}{\lambda_{\min}(\bar{\mathbf{X}}^T \bar{\mathbf{X}})} \text{ is the condition number.}$$

The pseudo-code of CG is available at the [Wikipedia](#).

Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient

Least Squares Regression

- Solve the optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \left\| \bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} \right\|_2^2$$

Tasks

Linear
Regression

Methods

Least Squares Regression

LASSO

Least Absolute Deviations

Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient

Solve Least Squares in Python

1. Load Data

```
from keras.datasets import boston_housing

(x_train, y_train), (x_test, y_test) = boston_housing.load_data()

print('shape of x_train: ' + str(x_train.shape))
print('shape of x_test: ' + str(x_test.shape))
print('shape of y_train: ' + str(y_train.shape))
print('shape of y_test: ' + str(y_test.shape))
```

```
shape of x_train: (404, 13)
shape of x_test: (102, 13)
shape of y_train: (404,)
shape of y_test: (102,)
```

2. Add A Feature

```
import numpy

n, d = x_train.shape
xbar_train = numpy.concatenate((x_train, numpy.ones((n, 1))),
                                axis=1)

print('shape of x_train: ' + str(x_train.shape))
print('shape of xbar_train: ' + str(xbar_train.shape))

shape of x_train: (404, 13)
shape of xbar_train: (404, 14)
```


3. Solve the Least Squares

Analytical solution: $\bar{\mathbf{w}} = (\bar{\mathbf{X}}^T \bar{\mathbf{X}})^{-1} \bar{\mathbf{X}}^T \mathbf{y}$

```
xx = numpy.dot(xbar_train.T, xbar_train)
xx_inv = numpy.linalg.pinv(xx)
xy = numpy.dot(xbar_train.T, y_train)
w = numpy.dot(xx_inv, xy)
```

3. Solve the Least Squares

Analytical solution: $\bar{\mathbf{w}} = (\bar{\mathbf{X}}^T \bar{\mathbf{X}})^{-1} \bar{\mathbf{X}}^T \mathbf{y}$

```
xx = numpy.dot(xbar_train.T, xbar_train)
xx_inv = numpy.linalg.pinv(xx)
xy = numpy.dot(xbar_train.T, y_train)
w = numpy.dot(xx_inv, xy)
```

3. Solve the Least Squares

Analytical solution: $\bar{\mathbf{w}} = (\bar{\mathbf{X}}^T \bar{\mathbf{X}})^{-1} \bar{\mathbf{X}}^T \mathbf{y}$

```
xx = numpy.dot(xbar_train.T, xbar_train)
xx_inv = numpy.linalg.pinv(xx)
xy = numpy.dot(xbar_train.T, y_train)
w = numpy.dot(xx_inv, xy)
```

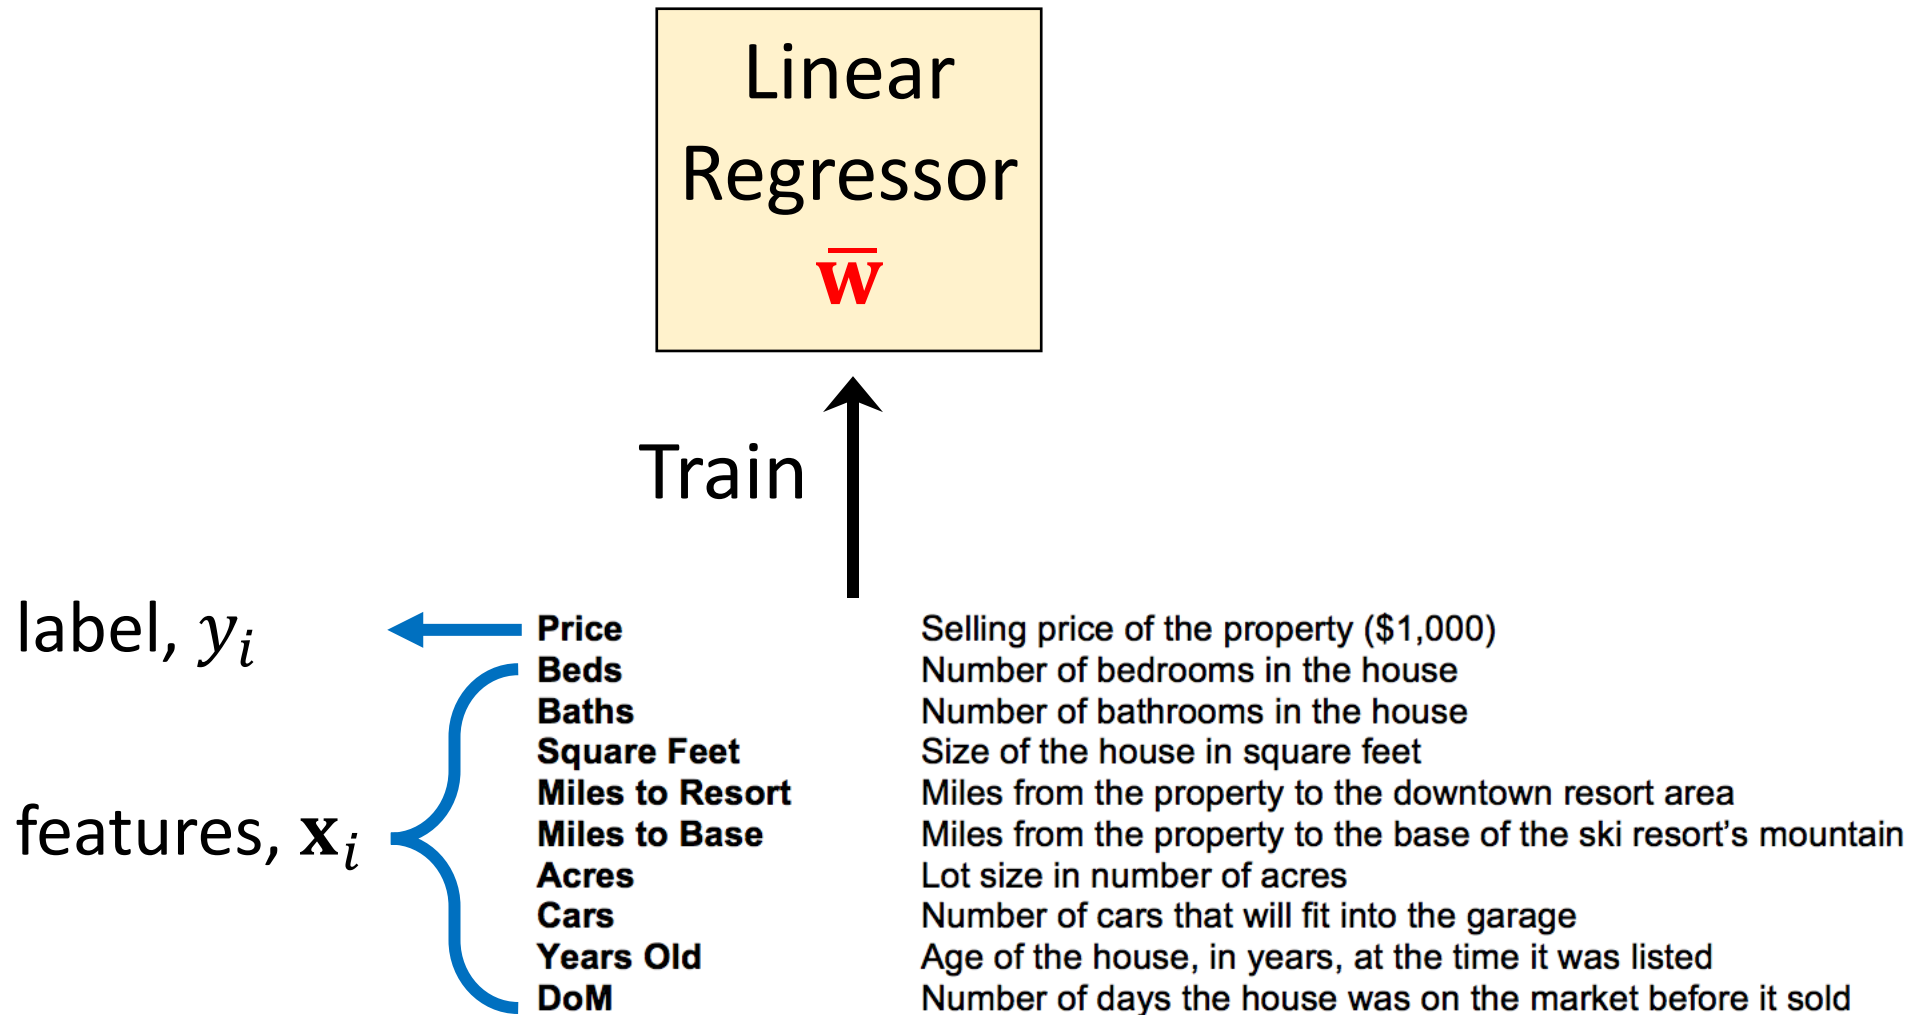
3. Solve the Least Squares

Training Mean Squared Error (MSE): $\frac{1}{n} \|\mathbf{y} - \bar{\mathbf{X}}\bar{\mathbf{w}}\|_2^2$

```
y_lsr = numpy.dot(xbar_train, w)
diff = y_lsr - y_train
mse = numpy.mean(diff * diff)
print('Train MSE: ' + str(mse))
```

Train MSE: 22.00480083834814

Linear Regression for Housing Price



Linear Regression for Housing Price



Features of a House, \mathbf{x}'
→ Extend it to $\bar{\mathbf{x}}'$

Linear
Regressor
 $\bar{\mathbf{w}}$

Predict

Price:
 $\bar{\mathbf{w}}^T \bar{\mathbf{x}}' = \500K

Train

label, y_i

features, \mathbf{x}_i

Price

Beds

Baths

Square Feet

Miles to Resort

Miles to Base

Acres

Cars

Years Old

DoM

Selling price of the property (\$1,000)

Number of bedrooms in the house

Number of bathrooms in the house

Size of the house in square feet

Miles from the property to the downtown resort area

Miles from the property to the base of the ski resort's mountain

Lot size in number of acres

Number of cars that will fit into the garage

Age of the house, in years, at the time it was listed

Number of days the house was on the market before it sold

4. Make Prediction for Test Samples

- Add a feature to the test feature matrix: $\mathbf{X}_{\text{test}} \rightarrow \bar{\mathbf{X}}_{\text{test}}$.
- Make prediction by: $\mathbf{y}_{\text{pred}} = \bar{\mathbf{X}}_{\text{test}} \bar{\mathbf{w}}$.

```
n_test, _ = x_test.shape
xbar_test = numpy.concatenate((x_test, numpy.ones((n_test, 1))), axis=1)
y_pred = numpy.dot(xbar_test, w)
```

4. Make Prediction for Test Samples

- Add a feature to the test feature matrix: $\mathbf{X}_{\text{test}} \rightarrow \bar{\mathbf{X}}_{\text{test}}$.
- Make prediction by: $\mathbf{y}_{\text{pred}} = \bar{\mathbf{X}}_{\text{test}} \bar{\mathbf{w}}$.
- MSE (test): $\frac{1}{n_{\text{test}}} \left\| \mathbf{y}_{\text{pred}} - \mathbf{y}_{\text{test}} \right\|_2^2$

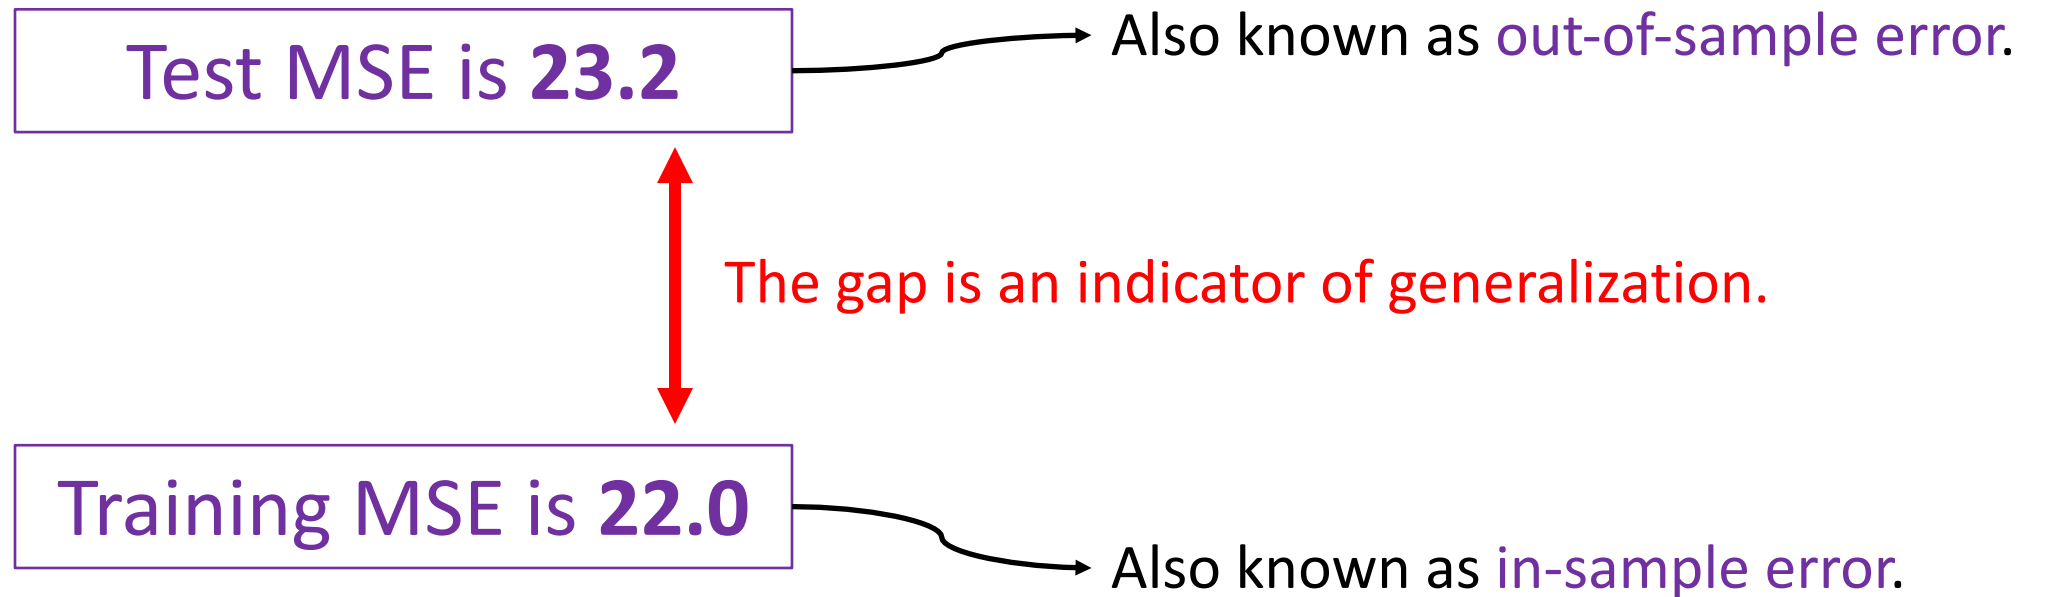
```
# mean squared error (testing)
```

```
diff = y_pred - y_test  
mse = numpy.mean(diff * diff)  
print('Test MSE: ' + str(mse))
```

Test MSE: 23.195599256409857

Training MSE is **22.0**

4. Make Prediction for Testing Samples



5. Compare with Baseline

Baseline:

- whatever the features are, the prediction is $\text{mean}(\mathbf{y})$.

```
y_mean = numpy.mean(y_train)

diff = y_pred - y_mean
mse = numpy.mean(diff * diff)
print('Test MSE: ' + str(mse))
```

Test MSE: 57.38297638530044

Test MSE of least
squares is **23.19**

Summary

- Linear regression problem.
- Least squares model.
- 3 algorithms for solving the model.
- Make predictions for never-seen-before test data.
- Evaluation of the model (training MSE and test MSE).
- Compare with baselines.