

# Linear Regression

Shusen Wang

## **Warm-up: Vector and Matrix**

# Vector and Matrix

Vector ( $n$ -dim)       $\mathbf{a} = \begin{bmatrix} a_1 \\ a_2 \\ \vdots \\ a_n \end{bmatrix}$


Matrix ( $n \times d$ )       $\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1d} \\ a_{21} & a_{22} & \cdots & a_{2d} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \cdots & a_{nd} \end{bmatrix}$

Row and columns       $\mathbf{A} = \begin{bmatrix} \mathbf{a}_{:1} & \mathbf{a}_{:2} & \cdots & \mathbf{a}_{:d} \end{bmatrix} = \begin{bmatrix} \mathbf{a}_{1:} \\ \mathbf{a}_{2:} \\ \vdots \\ \mathbf{a}_{n:} \end{bmatrix}$

# Vector Norms

- The  $\ell_p$  norm:  $\|\mathbf{x}\|_p := (\sum_i |x_i|^p)^{1/p}$ .
- The  $\ell_2$  norm:  $\|\mathbf{x}\|_2 = (\sum_i x_i^2)^{1/2}$  (the Euclidean norm).
- The  $\ell_1$  norm  $\|\mathbf{x}\|_1 = \sum_i |x_i|$ .
- The  $\ell_\infty$  norm is defined by  $\|\mathbf{x}\|_\infty = \max_i |x_i|$ .

# Vector Norms



- The  $\ell_2$ -distance (Euclidean distance):  
 $\|\mathbf{a} - \mathbf{b}\|_2$  (green line)
- The  $\ell_1$ -distance (Manhattan distance):  
 $\|\mathbf{a} - \mathbf{b}\|_1$  (red, blue, yellow lines)

# Transpose and Rank

Transpose:  $\begin{bmatrix} 6 & 4 & 24 \\ 1 & -9 & 8 \end{bmatrix}^T = \begin{bmatrix} 6 & 1 \\ 4 & -9 \\ 24 & 8 \end{bmatrix}$

**Square matrix:** a matrix with the same number of rows and columns.

**Symmetric:** a square matrix  $\mathbf{A}$  is symmetric if  $\mathbf{A}^T = \mathbf{A}$ .

**Rank:** the number of linearly independent rows (or columns).

**Full rank:** a square matrix is full rank if the rank equals to #columns.

# Eigenvalue Decomposition

- Let  $\mathbf{A}$  be any  $n \times n$  symmetric matrix.
- Eigenvalue decomposition:  $\mathbf{A} = \sum_{i=1}^n \lambda_i \mathbf{v}_i \mathbf{v}_i^T$ .
- Eigenvalues satisfy  $|\lambda_1| \geq |\lambda_2| \geq \dots \geq |\lambda_n|$ .
- Eigenvectors satisfy  $\mathbf{v}_i^T \mathbf{v}_j = 0$  for all  $i \neq j$ .
- $\mathbf{A}$  is full rank  $\iff$  all the eigenvalues are nonzero.

# **Warm-up: Optimization**



# Optimization: Basics

Optimization problem:  $\min_{\mathbf{w}} f(\mathbf{w}); \quad \text{s. t. } \mathbf{w} \in \mathcal{C}.$

- $\mathbf{w} = [w_1, \dots, w_d]$  : optimization variables
- $f : \mathbb{R}^d \mapsto \mathbb{R}$  : objective function
- $\mathcal{C}$  (a subset of  $\mathbb{R}^d$ ) : feasible set

# Optimization: Basics

Optimization problem:  $\min_{\mathbf{w}} f(\mathbf{w});$  s. t.  $\mathbf{w} \in \mathcal{C}$ .

- $\mathbf{w} = [w_1, \dots, w_d]$  : optimization variables
- $f : \mathbb{R}^d \mapsto \mathbb{R}$  : objective function
- $\mathcal{C}$  (a subset of  $\mathbb{R}^d$ ) : feasible set

  
Constraint

# Optimization: Basics

Optimization problem:  $\min_{\mathbf{w}} f(\mathbf{w}); \quad \text{s. t. } \mathbf{w} \in \mathcal{C}.$

- $\mathbf{w} = [w_1, \dots, w_d]$  : optimization variables
- $f : \mathbb{R}^d \mapsto \mathbb{R}$  : objective function
- $\mathcal{C}$  (a subset of  $\mathbb{R}^d$ ) : feasible set
  
- $\mathbf{w}^* = \operatorname{argmin}_{\mathbf{w} \in \mathcal{C}} f(\mathbf{w})$  is the optimal solution to the problem
  - $f(\mathbf{w}^*) \leq f(\mathbf{w})$  for all the vectors  $\mathbf{w}$  in the set  $\mathcal{C}$ .

# **Least Squares Regression**

# The Linear Regression Task


**Input:** vectors  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  and labels  $y_1, \dots, y_n \in \mathbb{R}$

**Output:** a vector  $\mathbf{w} \in \mathbb{R}^d$  and scalar  $b \in \mathbb{R}$  such that  $\mathbf{x}_i^T \mathbf{w} + b \approx y_i$ .

1-dim ( $d = 1$ ) example:

Solution:

$$y_i \approx 0.15 x_i + 5.0$$




# The Linear Regression Task

**Input:** vectors  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  and labels  $y_1, \dots, y_n \in \mathbb{R}$

**Output:** a vector  $\mathbf{w} \in \mathbb{R}^d$  and scalar  $b \in \mathbb{R}$  such that  $\mathbf{x}_i^T \mathbf{w} + b \approx y_i$ .

**Question** (regard training):  
how to compute  $\mathbf{w}$  and  $b$ ?



# The Linear Regression Task

**Input:** vectors  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  and labels  $y_1, \dots, y_n \in \mathbb{R}$

**Output:** a vector  $\mathbf{w} \in \mathbb{R}^d$  and scalar  $b \in \mathbb{R}$  such that  $\mathbf{x}_i^T \mathbf{w} + b \approx y_i$ .

**Method:** least squares regression.

- The optimization model:

$$\min_{\mathbf{w}, b} L(\mathbf{w}, b), \quad \text{where } L(\mathbf{w}, b) = \sum_{i=1}^n (\mathbf{x}_i^T \mathbf{w} + b - y_i)^2$$

# Least Squares Regression


- The optimization model:

$$\min_{\mathbf{w}, b} L(\mathbf{w}, b), \quad \text{where } L(\mathbf{w}, b) = \sum_{i=1}^n (\mathbf{x}_i^T \mathbf{w} + b - y_i)^2$$



# Least Squares Regression

- The optimization model:

$$\min_{\mathbf{w}, b} L(\mathbf{w}, b), \quad \text{where } L(\mathbf{w}, b) = \sum_{i=1}^n (\mathbf{x}_i^T \mathbf{w} + b - y_i)^2$$


Intercept

# Least Squares Regression

- The optimization model:

$$\min_{\mathbf{w}, b} L(\mathbf{w}, b), \quad \text{where } L(\mathbf{w}, b) = \sum_{i=1}^n (\mathbf{x}_i^T \mathbf{w} + b - y_i)^2$$




$$\min_{\mathbf{w} \in \mathbb{R}^{d+1}} \sum_{i=1}^n (\bar{\mathbf{x}}_i^T \bar{\mathbf{w}} - y_i)^2$$

- Define  $\bar{\mathbf{x}}_i = [\mathbf{x}_i; 1] \in \mathbb{R}^{d+1}$
- Define  $\bar{\mathbf{w}} = [\mathbf{w}, b] \in \mathbb{R}^{d+1}$
- $\rightarrow \mathbf{x}_i^T \mathbf{w} + b = \bar{\mathbf{x}}_i^T \bar{\mathbf{w}}$

# Least Squares Regression

- The optimization model:

$$\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \sum_{i=1}^n (\bar{\mathbf{x}}_i^T \bar{\mathbf{w}} - y_i)^2$$


Matrix form:  $\min_{\bar{\mathbf{w}} \in \mathbb{R}^{d+1}} \|\bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y}\|_2^2$

$$\bar{\mathbf{X}} = \begin{bmatrix} \mathbf{x}_1^T & 1 \\ \mathbf{x}_2^T & 1 \\ \vdots & \vdots \\ \mathbf{x}_n^T & 1 \end{bmatrix}$$

$$n \times (d+1)$$

$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

$$n \times 1$$

$$\bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y} = \begin{bmatrix} \bar{\mathbf{x}}_1^T \bar{\mathbf{w}} - y_1 \\ \bar{\mathbf{x}}_2^T \bar{\mathbf{w}} - y_2 \\ \vdots \\ \bar{\mathbf{x}}_n^T \bar{\mathbf{w}} - y_n \end{bmatrix}$$

$$n \times 1$$

# Least Squares Regression

- The optimization model:

$$\min_{\mathbf{w} \in \mathbb{R}^{d+1}} \|\bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y}\|_2^2$$

Tasks

Linear  
Regression

Methods

Least Squares Regression

LASSO

Least Absolute Deviations

Algorithms

?

# Least Squares Regression

- The optimization model:

$$\min_{\mathbf{w} \in \mathbb{R}^{d+1}} \|\bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y}\|_2^2$$

Tasks

Linear  
Regression

Methods

Least Squares Regression

LASSO

Least Absolute Deviations

Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient (CG)

# Least Squares Regression

- Solve the optimization model:

$$\min_{\mathbf{w} \in \mathbb{R}^{d+1}} \|\bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y}\|_2^2$$

Gradient:  $\frac{\partial \|\mathbf{X} \mathbf{w} - \mathbf{y}\|_2^2}{\partial \mathbf{w}} = 2(\bar{\mathbf{X}}^T \bar{\mathbf{X}} \bar{\mathbf{w}} - \bar{\mathbf{X}}^T \mathbf{y})$

## Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient

# Least Squares Regression

- Solve the optimization model:

$$\min_{\mathbf{w} \in \mathbb{R}^{d+1}} \|\bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y}\|_2^2$$

Gradient:  $\frac{\partial \|\bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y}\|_2^2}{\partial \bar{\mathbf{w}}} = 2(\bar{\mathbf{X}}^T \bar{\mathbf{X}} \bar{\mathbf{w}} - \bar{\mathbf{X}}^T \mathbf{y}) = \mathbf{0}$

1<sup>st</sup>-order optimal condition

**Algorithms**

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient

# Least Squares Regression

- Solve the optimization model:

$$\min_{\mathbf{w} \in \mathbb{R}^{d+1}} \|\bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y}\|_2^2$$

Gradient:  $\frac{\partial \|\bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y}\|_2^2}{\partial \bar{\mathbf{w}}} = 2(\bar{\mathbf{X}}^T \bar{\mathbf{X}} \bar{\mathbf{w}} - \bar{\mathbf{X}}^T \mathbf{y}) = \mathbf{0}$



Normal equation:  $\bar{\mathbf{X}}^T \bar{\mathbf{X}} \bar{\mathbf{w}} = \bar{\mathbf{X}}^T \mathbf{y}$



Assume  $\bar{\mathbf{X}}^T \bar{\mathbf{X}}$  is full rank.

Analytical solution:  $\bar{\mathbf{w}} = (\bar{\mathbf{X}}^T \bar{\mathbf{X}})^{-1} \bar{\mathbf{X}}^T \mathbf{y}$

## Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient



# Least Squares Regression

- Solve the optimization model:

$$\min_{\mathbf{w} \in \mathbb{R}^{d+1}} \|\bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y}\|_2^2$$

$$\text{Gradient: } \frac{\partial \|\mathbf{X} \mathbf{w} - \mathbf{y}\|_2^2}{\partial \mathbf{w}} = 2(\bar{\mathbf{X}}^T \bar{\mathbf{X}} \bar{\mathbf{w}} - \bar{\mathbf{X}}^T \mathbf{y}) = \mathbf{0}$$

Gradient descent repeats:

1. Compute gradient:  $\mathbf{g}_t = \bar{\mathbf{X}}^T \bar{\mathbf{X}} \bar{\mathbf{w}}_t - \bar{\mathbf{X}}^T \mathbf{y}$
2. Update:  $\bar{\mathbf{w}}_{t+1} = \bar{\mathbf{w}}_t - \alpha_t \mathbf{g}_t$

**Algorithms**

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient

# Least Squares Regression

- Solve the optimization model:

$$\min_{\mathbf{w} \in \mathbb{R}^{d+1}} \|\bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y}\|_2^2$$

Convergence: after  $O\left(\kappa \log \frac{1}{\epsilon}\right)$  iterations,

$$\|\bar{\mathbf{X}} (\bar{\mathbf{w}}_t - \bar{\mathbf{w}}^*)\|_2 \leq \epsilon \|\bar{\mathbf{X}} (\bar{\mathbf{w}}_0 - \bar{\mathbf{w}}^*)\|_2.$$

$$\kappa = \frac{\lambda_{\max}(\mathbf{X}^T \mathbf{X})}{\lambda_{\min}(\mathbf{X}^T \mathbf{X})} \text{ is the condition number.}$$

## Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient

# Least Squares Regression

- Solve the optimization model:

$$\min_{\mathbf{w} \in \mathbb{R}^{d+1}} \|\bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y}\|_2^2$$

Convergence: after  $O\left(\sqrt{\kappa} \log \frac{1}{\epsilon}\right)$  iterations,

$$\|\bar{\mathbf{X}} (\bar{\mathbf{w}}_t - \bar{\mathbf{w}}^*)\|_2 \leq \epsilon \|\bar{\mathbf{X}} (\bar{\mathbf{w}}_0 - \bar{\mathbf{w}}^*)\|_2.$$

$$\kappa = \frac{\lambda_{\max}(\mathbf{X}^T \mathbf{X})}{\lambda_{\min}(\mathbf{X}^T \mathbf{X})} \text{ is the condition number.}$$

The pseudo-code of CG is available at the [Wikipedia](#).

## Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient

# Least Squares Regression

- Solve the optimization model:

$$\min_{\mathbf{w} \in \mathbb{R}^{d+1}} \|\bar{\mathbf{X}} \bar{\mathbf{w}} - \mathbf{y}\|_2^2$$

Tasks

Linear  
Regression

Methods

Least Squares Regression

LASSO

Least Absolute Deviations

Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient

# **Implement Least Squares in Python**

# 1. Load Data

```
from keras.datasets import boston_housing

(x_train, y_train), (x_test, y_test) = boston_housing.load_data()

print('shape of x_train: ' + str(x_train.shape))
print('shape of x_test: ' + str(x_test.shape))
print('shape of y_train: ' + str(y_train.shape))
print('shape of y_test: ' + str(y_test.shape))
```

```
shape of x_train: (404, 13)
shape of x_test: (102, 13)
shape of y_train: (404,)
shape of y_test: (102,)
```

## 2. Add A Feature

```
import numpy

n, d = x_train.shape
xbar_train = numpy.concatenate((x_train, numpy.ones((n, 1))),
                                axis=1)

print('shape of x_train: ' + str(x_train.shape))
print('shape of xbar_train: ' + str(xbar_train.shape))
```

---

```
shape of x_train: (404, 13)
shape of xbar_train: (404, 14)
```

### 3. Solve the Least Squares

Analytical solution:  $\bar{\mathbf{w}} = (\bar{\mathbf{X}}^T \bar{\mathbf{X}})^{-1} \bar{\mathbf{X}}^T \mathbf{y}$

```
xx = numpy.dot(xbar_train.T, xbar_train)
xx_inv = numpy.linalg.pinv(xx)
xy = numpy.dot(xbar_train.T, y_train)
w = numpy.dot(xx_inv, xy)
```



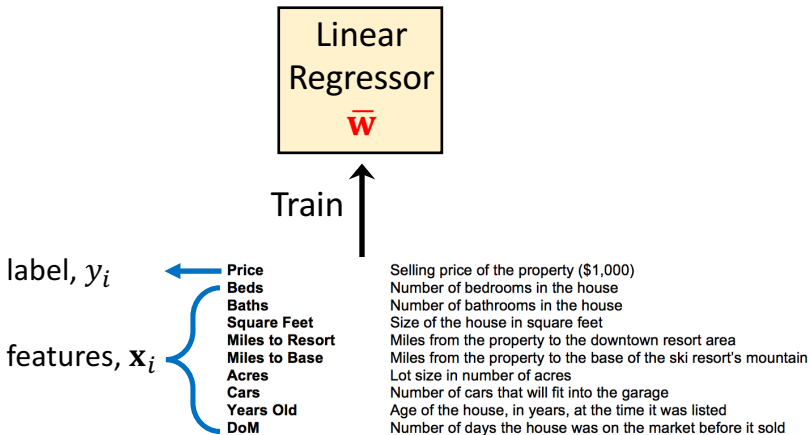
### 3. Solve the Least Squares

Mean squared error (training):  $\frac{1}{n} ||\mathbf{y} - \bar{\mathbf{X}}\bar{\mathbf{w}}||_2^2$

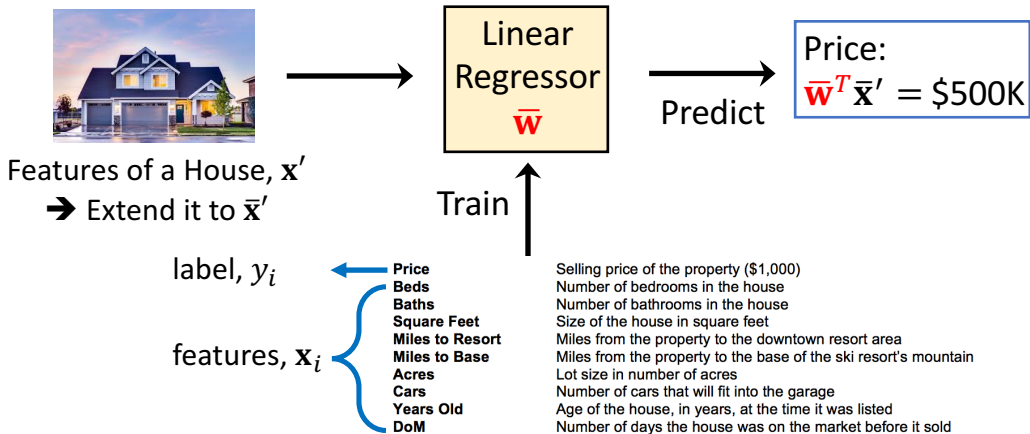
```
y_lsr = numpy.dot(xbar_train, w)
diff = y_lsr - y_train
mse = numpy.mean(diff * diff)
print('Train MSE: ' + str(mse))
```

Train MSE: 22.00480083834814

# Linear Regression for Housing Price



# Linear Regression for Housing Price



## 4. Make Prediction for Testing Samples

- Add a feature to the testing samples:  $\mathbf{X}_{\text{test}} \rightarrow \bar{\mathbf{X}}_{\text{test}}$ .
- Make prediction by:  $\mathbf{y}_{\text{pred}} = \bar{\mathbf{X}}_{\text{test}} \bar{\mathbf{w}}$ .

```
n_test, _ = x_test.shape
xbar_test = numpy.concatenate((x_test, numpy.ones((n_test, 1))), axis=1)
y_pred = numpy.dot(xbar_test, w)
```

## 4. Make Prediction for Testing Samples

- Add a feature to the testing samples:  $\mathbf{X}_{\text{test}} \rightarrow \bar{\mathbf{X}}_{\text{test}}$ .
- Make prediction by:  $\mathbf{y}_{\text{pred}} = \bar{\mathbf{X}}_{\text{test}} \bar{\mathbf{w}}$ .
- MSE (testing):  $\frac{1}{n_{\text{test}}} \left\| \mathbf{y}_{\text{pred}} - \mathbf{y}_{\text{test}} \right\|_2^2$

```
# mean squared error (testing)
```

```
diff = y_pred - y_test  
mse = numpy.mean(diff * diff)  
print('Test MSE: ' + str(mse))
```

Test MSE: 23.195599256409857

Training MSE is **22.0**

## 5. Compare with Baseline

Baseline:

- whatever the features are, the prediction is  $\text{mean}(\mathbf{y})$ .

```
y_mean = numpy.mean(y_train)

diff = y_pred - y_mean
mse = numpy.mean(diff * diff)
print('Test MSE: ' + str(mse))
```

Test MSE: 57.38297638530044

Test MSE of least  
squares is **23.19**