

# Regression

Shusen Wang

## **Warm-up: Least Squares Regression**

# Linear Regression (Task)


**Input:** vectors  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  and labels  $y_1, \dots, y_n \in \mathbb{R}$

**Output:** a vector  $\mathbf{w} \in \mathbb{R}^d$  and scalar  $b \in \mathbb{R}$  such that  $\mathbf{x}_i^T \mathbf{w} + b \approx y_i$ .

Tasks

Linear  
Regression

Inherently assume  $y_i$  is a  
linear function of  $\mathbf{x}_i$ .



# Least Squares Regression (Method)

**Input:** vectors  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  and labels  $y_1, \dots, y_n \in \mathbb{R}$

1. Add one dimension to  $\mathbf{x} \in \mathbb{R}^d$ :  $\bar{\mathbf{x}}_j = [\mathbf{x}_j; 1] \in \mathbb{R}^{d+1}$ .
2. Solve least squares regression:  $\min_{\mathbf{w} \in \mathbb{R}^{d+1}} \|\bar{\mathbf{X}} \mathbf{w} - \mathbf{y}\|_2^2$ .

Tasks

Methods

Linear  
Regression

Least Squares Regression

# Least Squares Regression (Method)

**Input:** vectors  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  and labels  $y_1, \dots, y_n \in \mathbb{R}$

1. Add one dimension to  $\mathbf{x} \in \mathbb{R}^d$ :  $\bar{\mathbf{x}}_j = [\mathbf{x}_j; 1] \in \mathbb{R}^{d+1}$ .
2. Solve least squares regression:  $\min_{\mathbf{w} \in \mathbb{R}^{d+1}} \|\bar{\mathbf{X}} \mathbf{w} - \mathbf{y}\|_2^2$ .

Tasks

Linear  
Regression

Methods

Least Squares Regression

Algorithms

Analytical Solution

Gradient Descent (GD)

Conjugate Gradient

# **Polynomial Regression**

# The Regression Task

**Input:** vectors  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  and labels  $y_1, \dots, y_n \in \mathbb{R}$ .

**Output:** a function  $f: \mathbb{R}^d \mapsto \mathbb{R}$  such that  $f(\mathbf{x}) \approx y$ .

**Question:**  $f$  is unknown! So how to learn  $f$ ?

**Answer:** polynomial approximation;  $f$  is a polynomial function.

# Polynomial Regression

**Input:** scalars  $x_1, \dots, x_n \in \mathbb{R}$  and labels  $y_1, \dots, y_n \in \mathbb{R}$ .

**Output:** a function  $f: \mathbb{R} \mapsto \mathbb{R}$  such that  $f(x) \approx y$ .

**One-dimensional example:**  $f(x) = w_0 + w_1x + w_2x^2 + \dots + w_px^p$ .

**Polynomial regression:**

1. Define a feature map  $\Phi(x) = [1, x, x^2, x^3, \dots, x^p]$ .
2. For  $j = 1$  to  $n$ , do the mapping  $x_j \mapsto \Phi(x_j)$ .
  - Let  $\Phi = [\Phi(x_1); \dots, \Phi(x_n)]^T \in \mathbb{R}^{n \times (p+1)}$
3. Solve the least squares regression  $\min_{\mathbf{w} \in \mathbb{R}^{p+1}} \|\Phi \mathbf{w} - \mathbf{y}\|_2^2$ .



# Polynomial Regression

**Input:** vectors  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^2$  and labels  $y_1, \dots, y_n \in \mathbb{R}$ .

**Output:** a function  $f: \mathbb{R}^2 \mapsto \mathbb{R}$  such that  $f(\mathbf{x}_i) \approx y_i$ .

**Two-dimensional example:** how to do feature mapping?

**Polynomial features:**

$$\Phi(\mathbf{x}) = [1, \quad x_1, \quad x_2, \quad x_1^2, \quad x_2^2, \quad x_1x_2, \quad x_1^3, \quad x_2^3, \quad x_1x_2^2, \quad x_1^2x_2].$$

0-order      1<sup>st</sup>-order      2<sup>nd</sup>-order      3<sup>rd</sup>-order

# Polynomial Regression

```
import numpy
X = numpy.arange(6).reshape(3, 2)
print('X = ')
print(X)
```

```
X =
[[0 1]
 [2 3]
 [4 5]]
```

```
from sklearn.preprocessing import PolynomialFeatures
poly = PolynomialFeatures(degree=3)
Phi = poly.fit_transform(X)
print('Phi = ')
print(Phi)
```

```
Phi =
[[ 1.  0.  1.  0.  0.  1.  0.  0.  0.  1.]
 [ 1.  2.  3.  4.  6.  9.  8. 12. 18. 27.]
 [ 1.  4.  5. 16. 20. 25. 64. 80. 100. 125.]]
```

0-order

1<sup>st</sup>-order

2<sup>nd</sup>-order

3<sup>rd</sup>-order


# Polynomial Regression

- $\mathbf{x}$ :  $d$ -dimensional
- $\Phi(\mathbf{x})$ :  $p$ -degree polynomial
- The dimension of  $\Phi(\mathbf{x})$  is  $O(d^p)$

# Polynomial Regression

**Input:** vectors  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  and labels  $y_1, \dots, y_n \in \mathbb{R}$ .

**Output:** a function  $f: \mathbb{R}^d \mapsto \mathbb{R}$  such that  $f(\mathbf{x}_i) \approx y_i$ .



# **Training, Testing, and Overfitting**

# Polynomial Regression: Training

**Input:** vectors  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  and labels  $y_1, \dots, y_n \in \mathbb{R}$ .


**Feature map:**  $\phi(\mathbf{x}) = \bigotimes^p \bar{\mathbf{x}}$ . Its dimension is  $O(d^p)$ .

**Least squares:**  $\min_{\mathbf{w}} \|\Phi \mathbf{w} - \mathbf{y}\|_2^2$ .



**Question:** what will happen as  $p$  grows?

1. For sufficiently large  $p$ , the dimension of the feature  $\phi(\mathbf{x})$  exceeds  $n$ .
2. Then you can find  $\mathbf{w}$  such that  $\Phi \mathbf{w} = \mathbf{y}$ .

# Polynomial Regression: Training



**Underfitting**



**Overfitting**

# Regression: Testing

**Train:**

**Input:** vectors  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  and labels  $y_1, \dots, y_n \in \mathbb{R}$ .

**Output:** a function  $f: \mathbb{R}^d \mapsto \mathbb{R}$  such that  $f(\mathbf{x}_i) \approx y_i$ .


**Test:**

**Input:** a feature vectors  $\mathbf{x}' \in \mathbb{R}^d$ .


**Input:** predict its label by  $f(\mathbf{x}')$ .




# Polynomial Regression: Testing



**BAD**




**GOOD**



**BAD**


# Training and Testing



# Training and Test

**Question:** for the polynomial regression model, how to determine the degree  $p$ ?

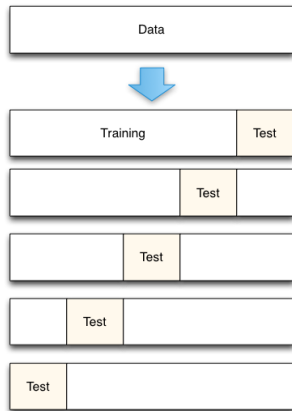
**Answer:** the degree  $p$  leads to the smallest test error.



# **$k$ -Fold Cross-Validation**

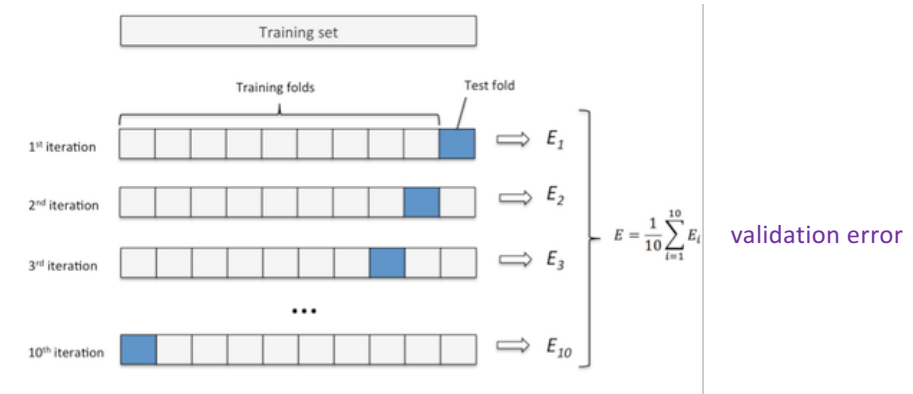
# $k$ -Fold Cross-Validation

1. Propose a grid of parameters, e.g.  $p \in \{1, 2, 3, 4\}$ .
2. Randomly partition the training samples to  $k$  parts.
  - $k - 1$  parts for training.
  - One part for test.
3. Compute the averaged test errors of the  $k$  repeats.
  - The average is called the **validation error**.
4. Choose the parameter  $p$  that leads to the smallest **validation error**.



Example: 5-fold cross-validation

# Example: 10-Fold Cross-Validation



# Example: 10-Fold Cross-Validation


parameter	validation error
p=1	23.19
p=2	21.00
p=3	18.54
p=4	24.36

**Remark:** cross-validation is performed on the training data.

# **Real-World Machine Learning Competition**




# The Available Data


	Training	Public	Private
Labels:	$y$	unknown	unknown
Features:	$X$	$X_{\text{public}}$	$X_{\text{private}}$
			
		Test Data	

The public and private are mixed;  
Participants cannot distinguish them.

# Train A Model



# Prediction



# Submission to Leaderboard

	Training	Public	Private
Labels:	$y$	unknown	unknown
Features:	$X$	$X_{\text{public}}$	$X_{\text{private}}$

Question: why two leaderboards?

