

Neural Machine Translation

Shusen Wang

Sequence-to-Sequence Model (Seq2Seq)

English

German

"do you agree" => **[Seq2Seq]** => "bist du einverstanden"

"go to sleep" => **[Seq2Seq]** => "gehen Sie schlafen"

"We will fight" => **[Seq2Seq]** => "Wir werden kämpfen"

Machine Translation Data

Datasets

- Tab-delimited Bilingual Sentence Pairs: <http://www.manythings.org/anki/>

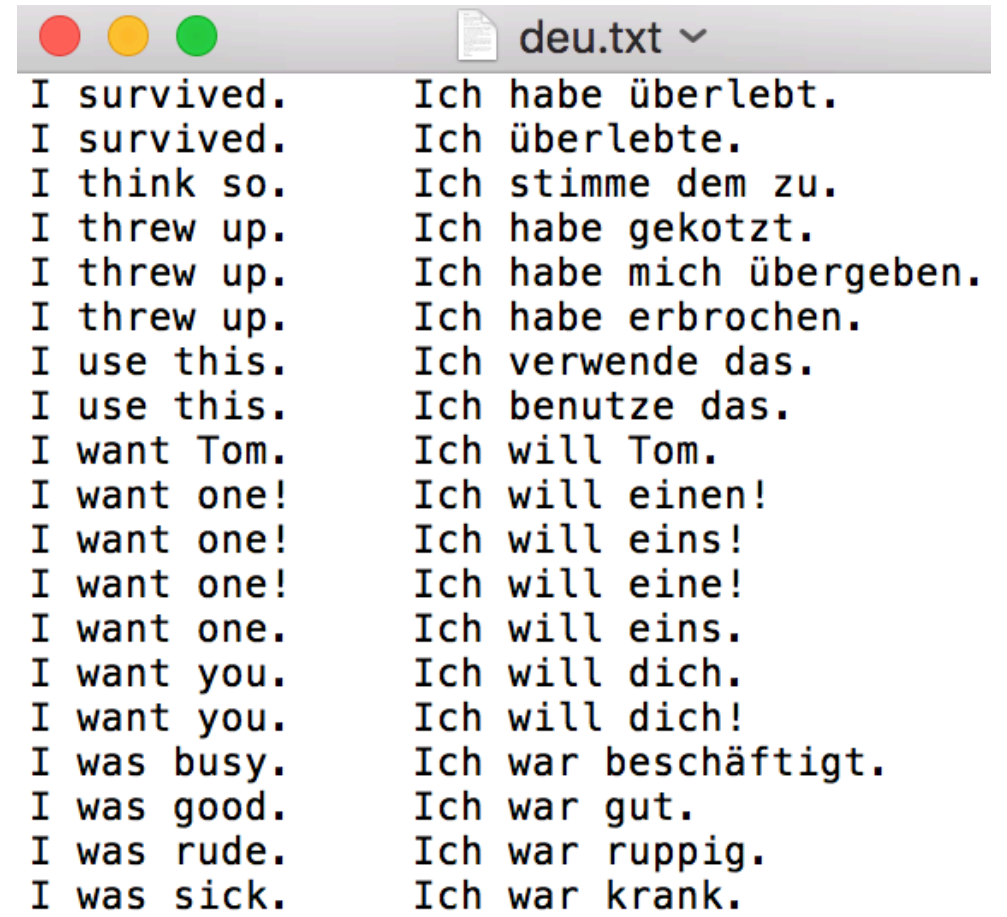


A sample of the German-English bilingual sentence pairs from the 'deu-eng.zip' dataset. The text is displayed in a window titled 'deu.txt'. Each line contains an English sentence followed by its German translation.

I survived.	Ich habe überlebt.
I survived.	Ich überlebte.
I think so.	Ich stimme dem zu.
I threw up.	Ich habe gekotzt.
I threw up.	Ich habe mich übergeben.
I threw up.	Ich habe erbrochen.
I use this.	Ich verwende das.
I use this.	Ich benutze das.
I want Tom.	Ich will Tom.
I want one!	Ich will einen!
I want one!	Ich will eins!
I want one!	Ich will eine!
I want one.	Ich will eins.
I want you.	Ich will dich.
I want you.	Ich will dich!
I was busy.	Ich war beschäftigt.
I was good.	Ich war gut.
I was rude.	Ich war ruppig.
I was sick.	Ich war krank.

1. Processing the Data

- **Preprocessing:** to lower case, remove punctuation, remove non-printable chars.



I survived.	Ich habe überlebt.
I survived.	Ich überlebte.
I think so.	Ich stimme dem zu.
I threw up.	Ich habe gekotzt.
I threw up.	Ich habe mich übergeben.
I threw up.	Ich habe erbrochen.
I use this.	Ich verwende das.
I use this.	Ich benutze das.
I want Tom.	Ich will Tom.
I want one!	Ich will einen!
I want one!	Ich will eins!
I want one!	Ich will eine!
I want one.	Ich will eins.
I want you.	Ich will dich.
I want you.	Ich will dich!
I was busy.	Ich war beschäftigt.
I was good.	Ich war gut.
I was rude.	Ich war ruppig.
I was sick.	Ich war krank.

2. Tokenization

- `input_texts` => [**Eng_Tokenizer**] => `input_tokens`
- `target_texts` => [**Deu_Tokenizer**] => `target_tokens`



Use 2 different tokenizers for the 2 languages

2. Tokenization

- `input_texts` => [**Eng_Tokenizer**] => `input_tokens`
- `target_texts` => [**Deu_Tokenizer**] => `target_tokens`

Eng_Tokenizer

- `"I_am_okay."` => `['i', '_', 'a', 'm', ..., 'a', 'y']`

Deu_Tokenizer

- `"Es geht mir gut"` => `['e', 's', '_', ..., 'u', 't']`

2. Tokenization

Question: Why 2 different tokenizers?

Answer: In the **char-level**, languages have different **alphabet/chars**.

- English: A a, B b, C c ..., Z z. (26 letters × 2).
- German: 26 letters, 3 umlauts (Ä, Ö, Ü) and one ligature (ß).
- Greek: Α α, Β β, Γ γ, Δ δ, ..., Ω ω. (24 letters × 2).
- Chinese: 金 木 水 火 土 ... 赵 钱 孙 李 (a few thousands characters).
- Japanese: あ い う え お ... (46 Hiragana, 46 Karagana, hundreds 漢字).

2. Tokenization

Question: Why 2 different tokenizers?

Answer: In the **word-level**, languages have different **vocabulary**.

- English:

Machine learning is a generic term for the artificial generation of knowledge from experience: An artificial system learns from examples and can generalize these after completion of the learning phase.

- Deutsche:



Maschinelles Lernen ist ein Oberbegriff für die künstliche Generierung von Wissen aus Erfahrung: Ein künstliches System lernt aus Beispielen und kann diese nach Beendigung der Lernphase verallgemeinern.

3. Encoding

Eng_Dictionary

- 'a' => 1
- 'b' => 2
- 'c' => 3
- 'd' => 4
- ...
- 'z' => 26
- '_' => 27

Deu_Dictionary

- '\t' => 1  start sign
- '\n' => 2  stop sign
- 'a' => 3
- 'b' => 4
- 'c' => 5
- 'd' => 6
- ...
- 'z' => 28
- '_' => 29

3. Encoding

"I_am_okay."

Eng_Tokenizer

['i', '_', 'a', 'm', '_', 'o', 'k', 'a', 'y']

Encoding using Eng_Dictionary

[9, 27, 1, 13, 27, 15, 11, 1, 25]

3. Encoding

"Es geht mir gut"

Deu_Tokenizer

['e', 's', '_', 'g', 'e', ..., 'g', 'u', 't']

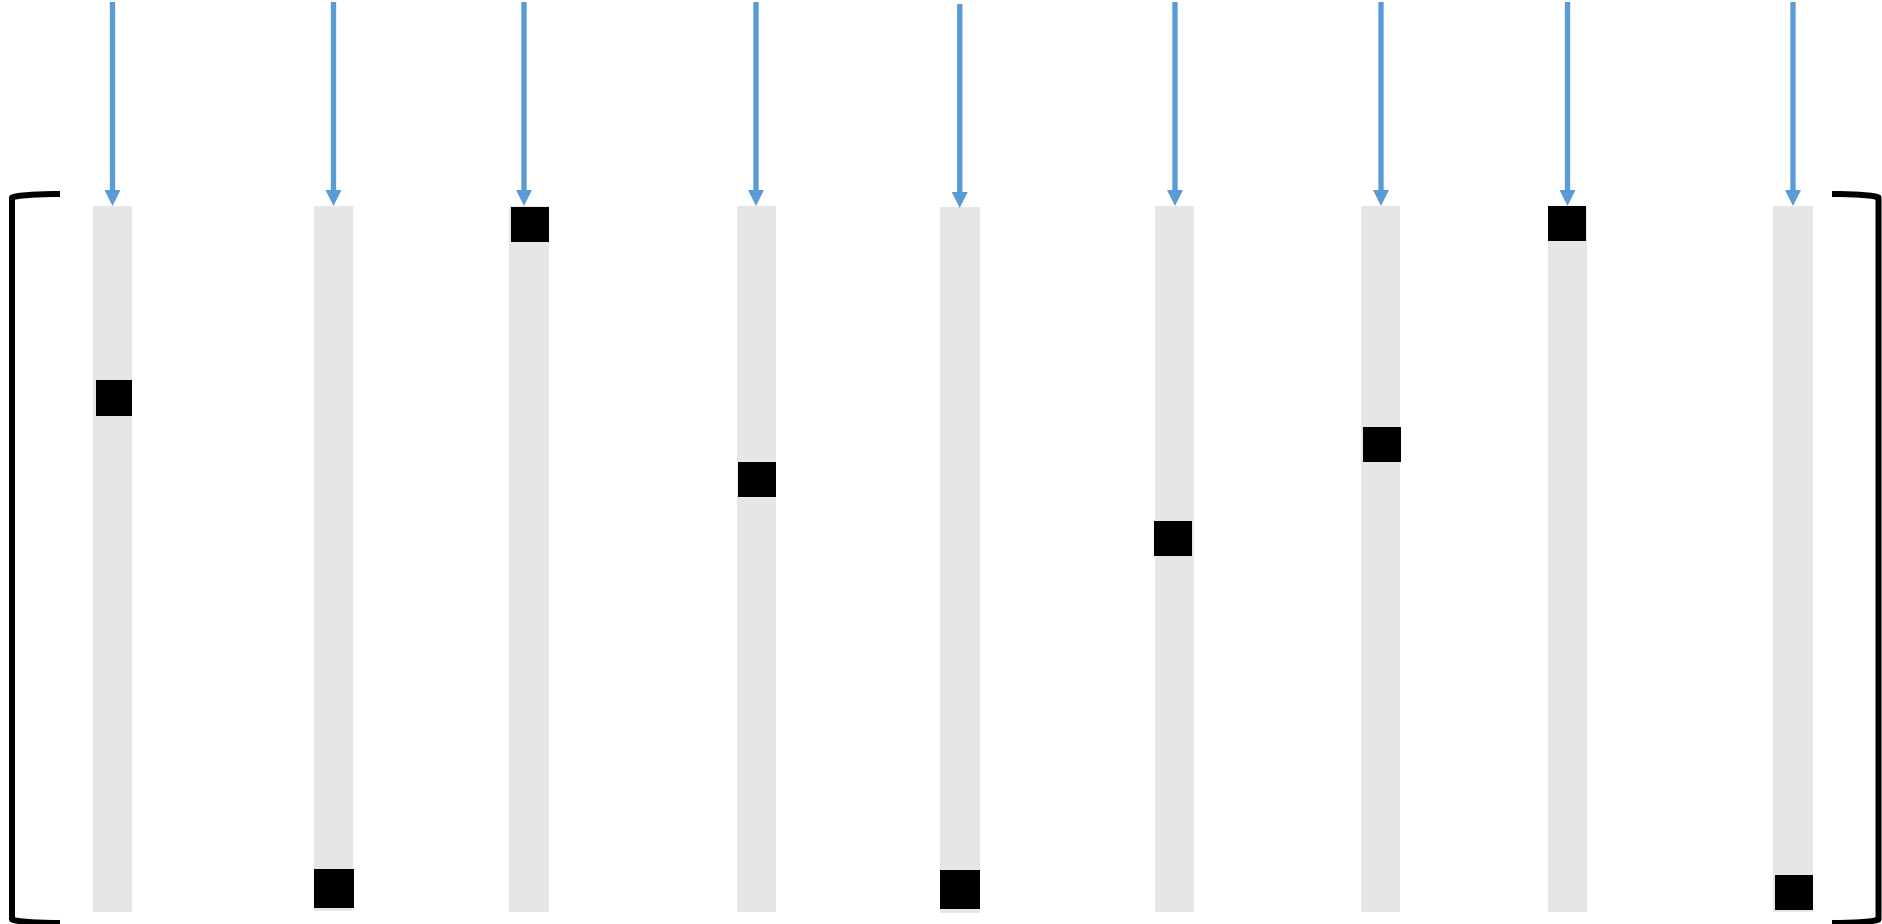
Encoding using Deu_Dictionary

[7, 21, 29, 9, 7, ..., 9, 23, 22]

4. Embedding (One-Hot Encoding)

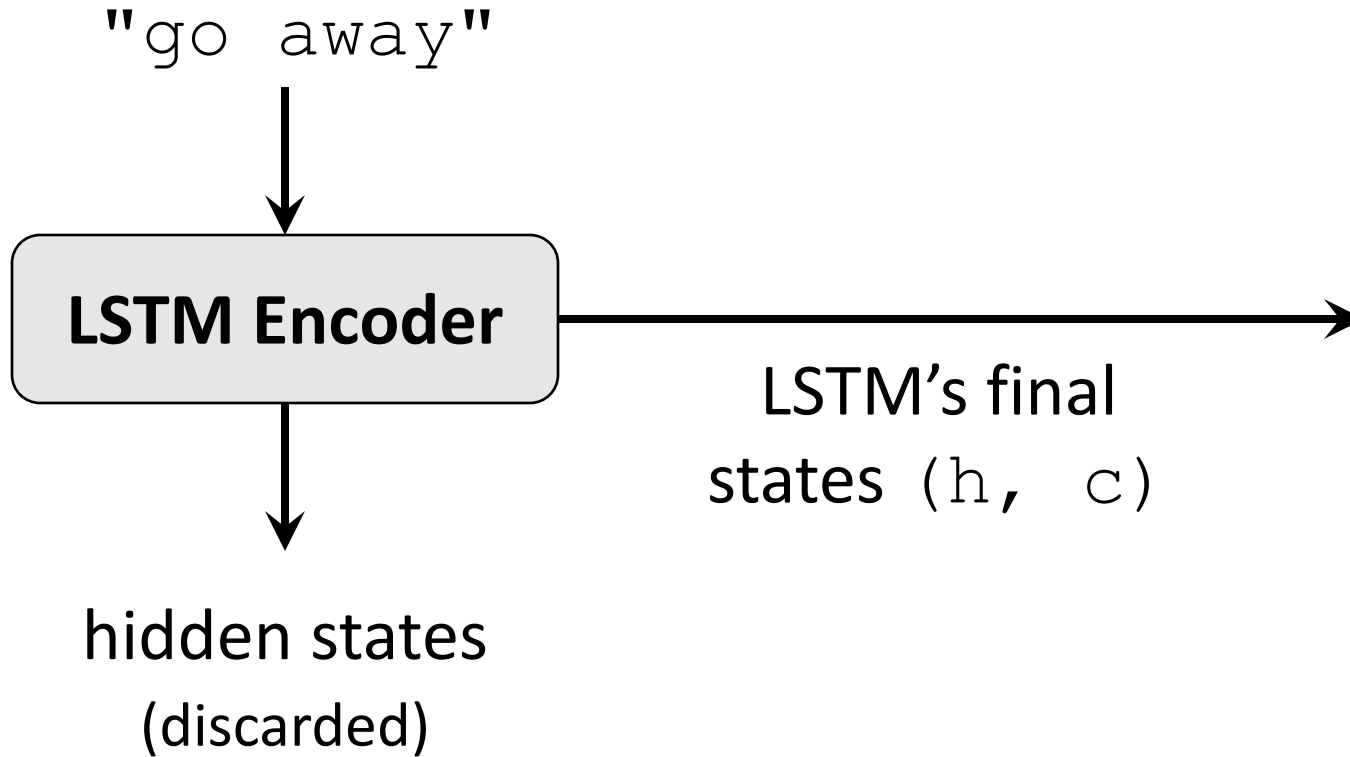
`['i', '_', 'a', 'm', '_', 'o', 'k', 'a', 'y']`

`[9, 27, 1, 13, 27, 15, 11, 1, 25]`



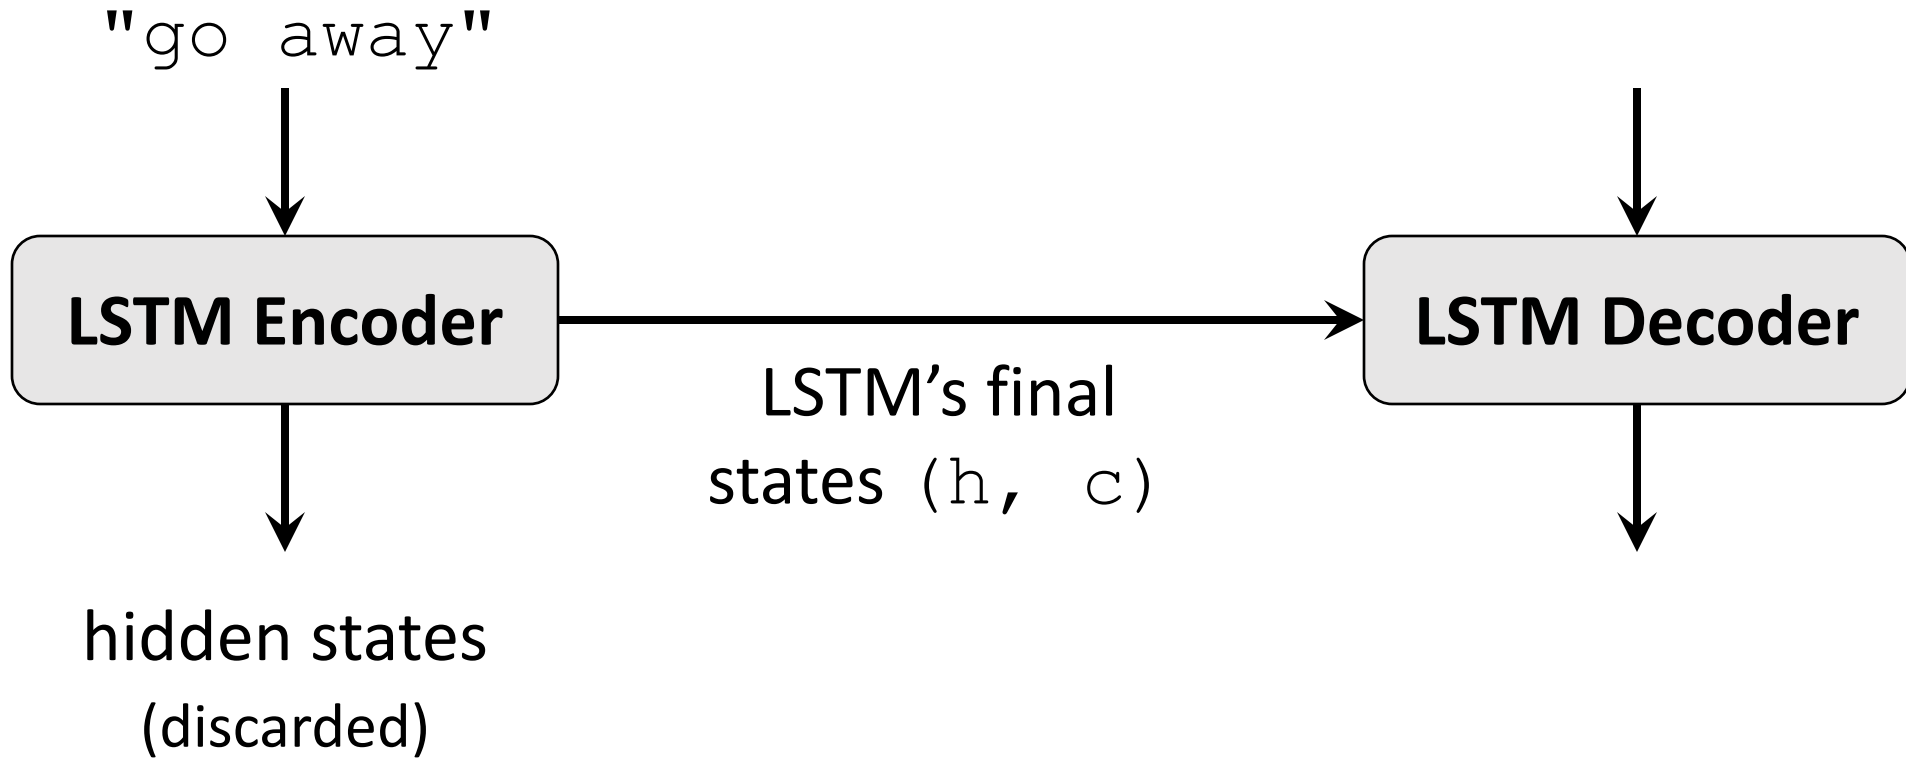
Seq2Seq Model

Training of the Seq2Seq Model

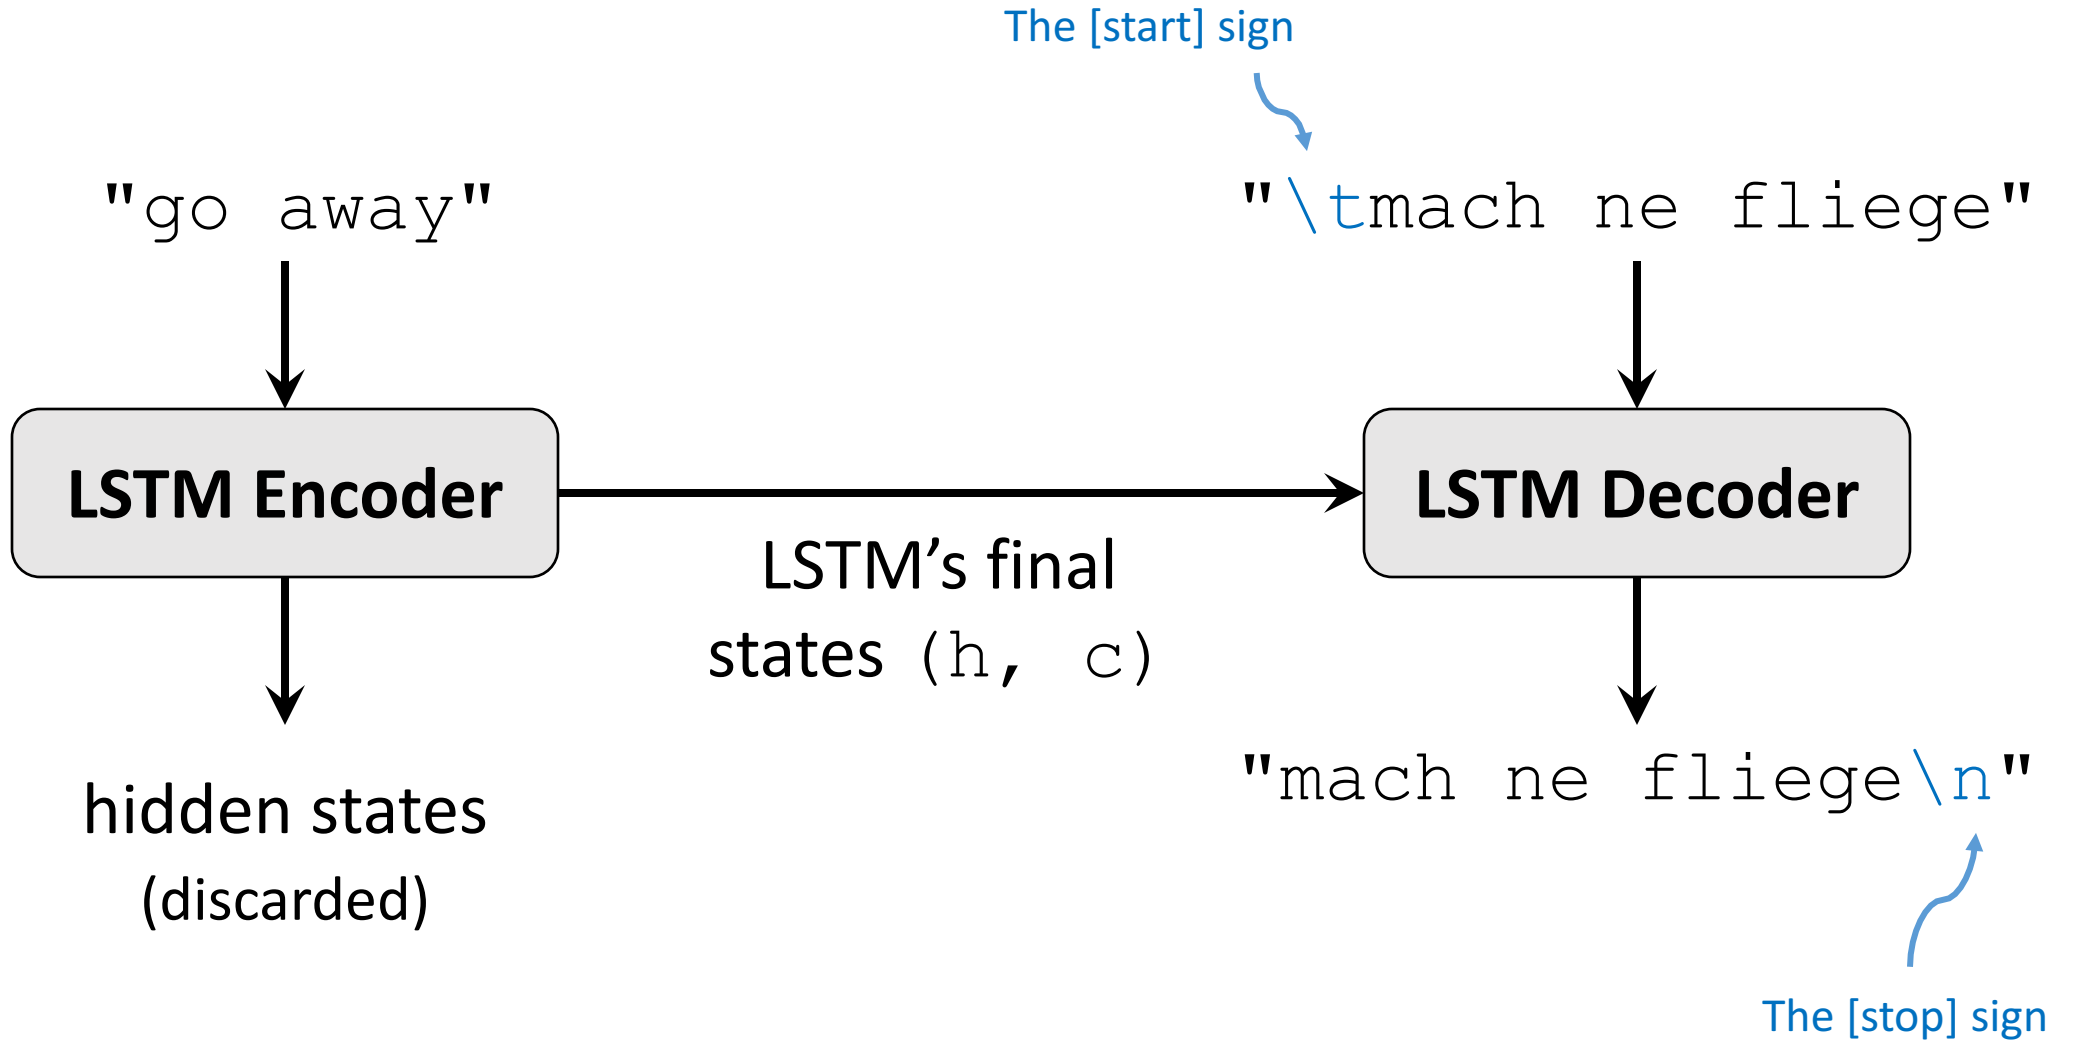


Training of the Seq2Seq Model

The [start] sign



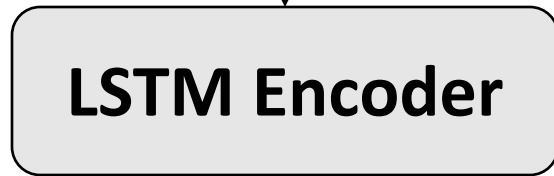
Training of the Seq2Seq Model



Training of the Seq2Seq Model

encoder_input[i]

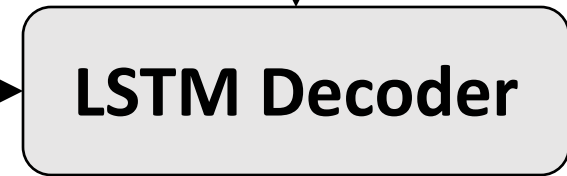
"go away"



hidden states
(discarded)

decoder_input[i]

"\tmach ne fliege"



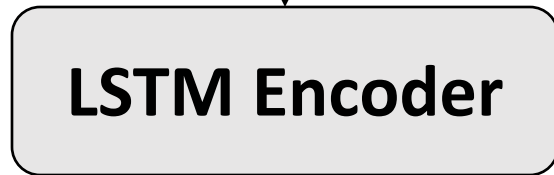
LSTM's final
states (h, c)

"mach ne fliege\n"

Training of the Seq2Seq Model

`encoder_input[i]`

"go away"

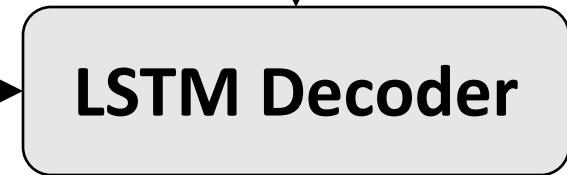


hidden states
(discarded)

LSTM's final
states (h , c)

`decoder_input[i]`

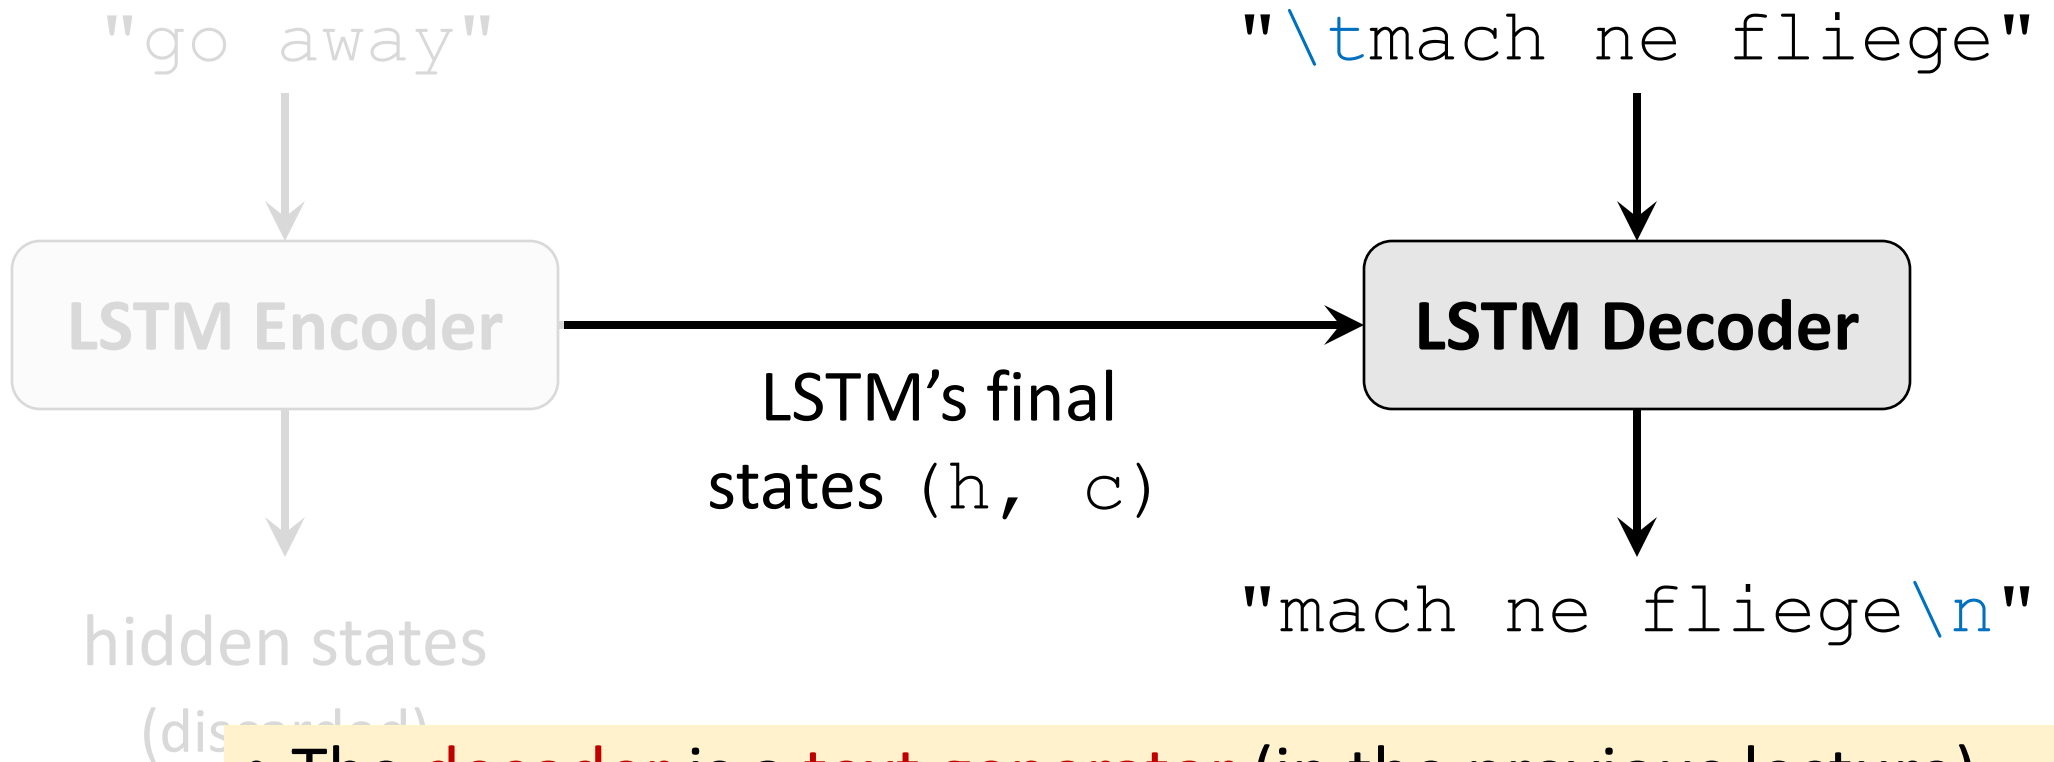
"\tmach ne fliege"



"mach ne fliege\n"

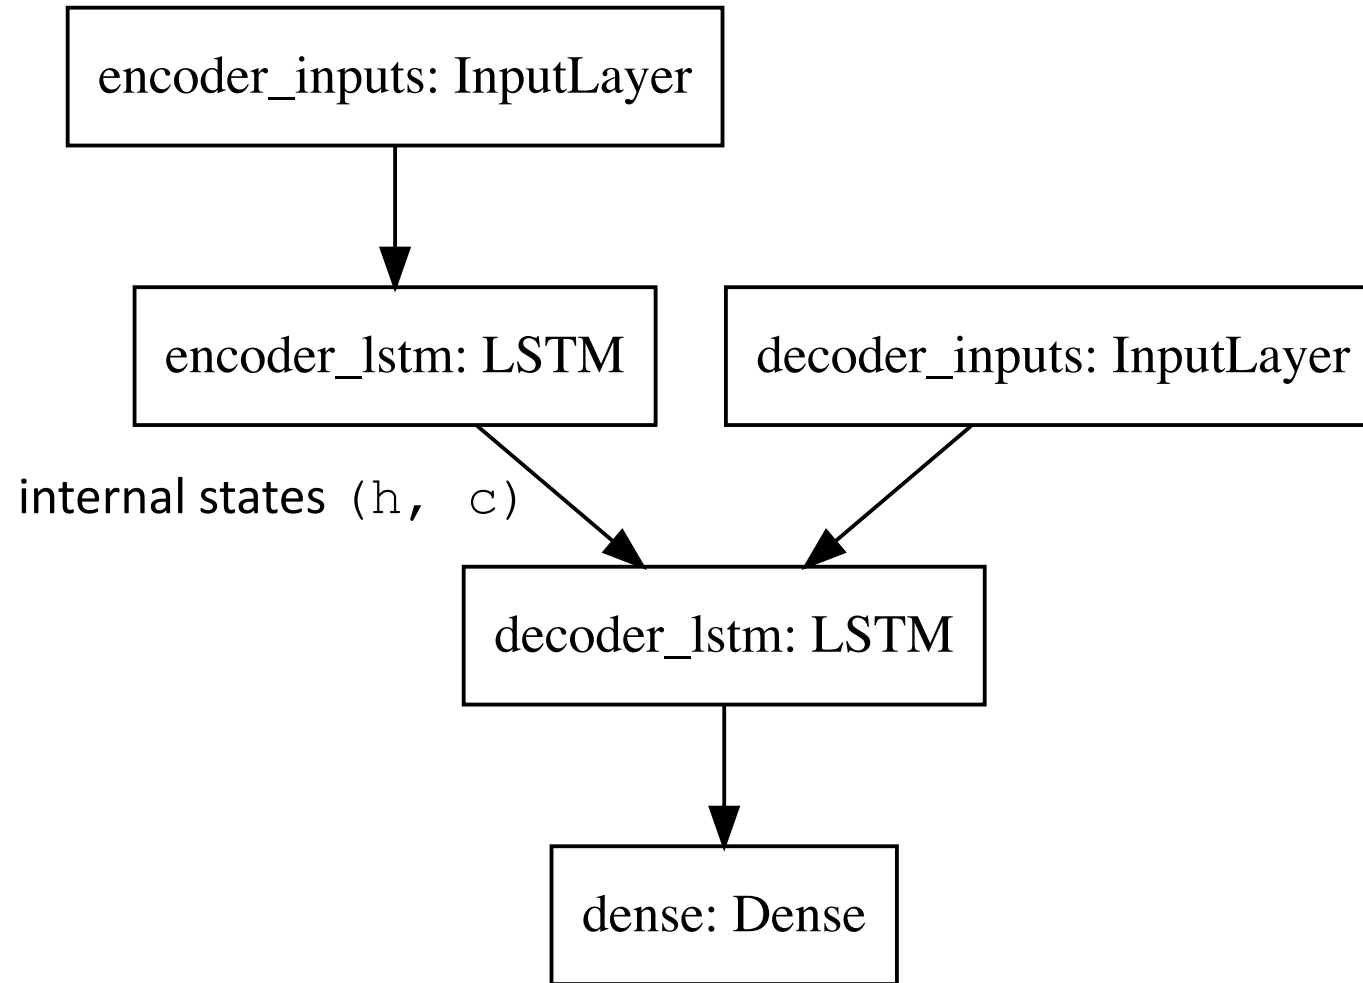
`decoder_target[i]`
(left shift of `decoder_input[i]`)

Training of the Seq2Seq Model



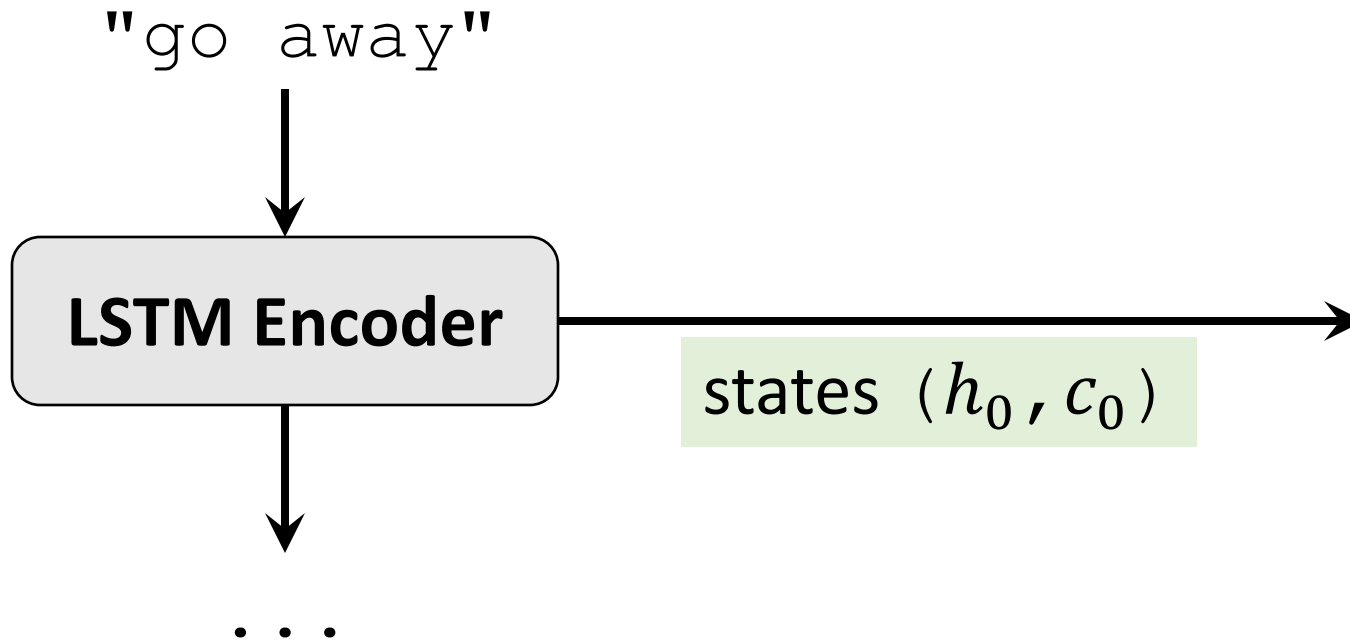
- The **decoder** is a **text generator** (in the previous lecture).
- **Difference** from the simple **text generator**: the **initial states** are determined by the **encoder**.

Seq2Seq Model in Keras

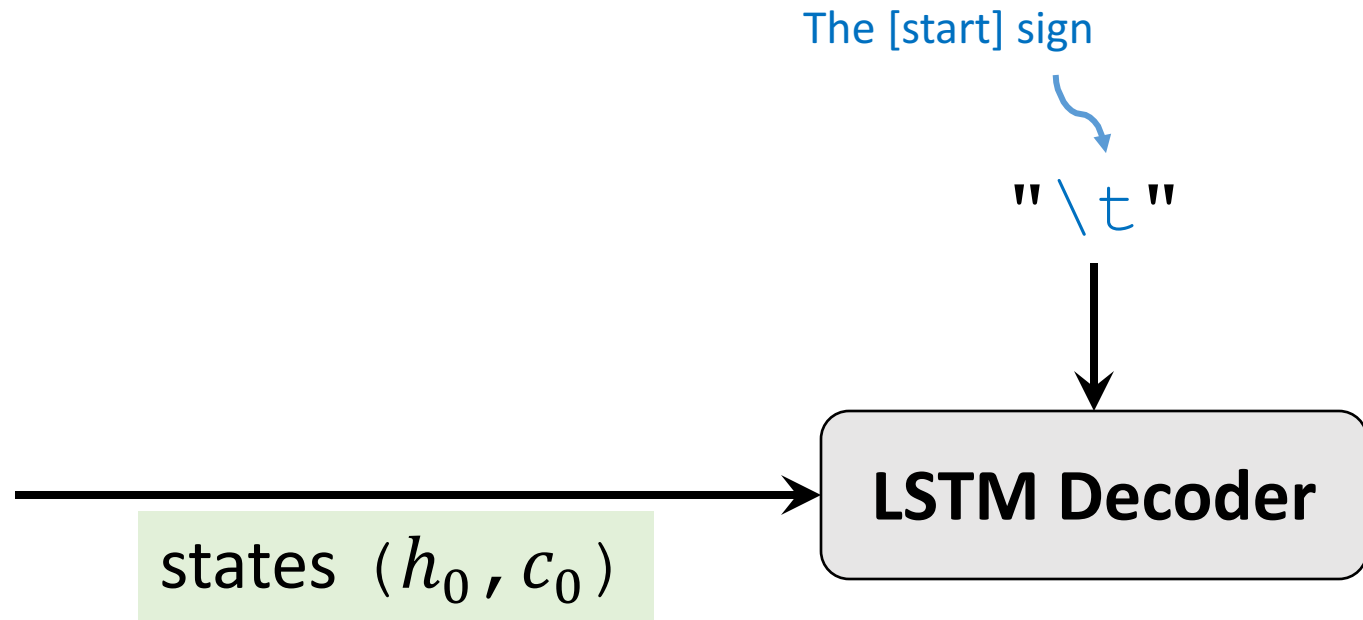


Inference Using the Seq2Seq Model

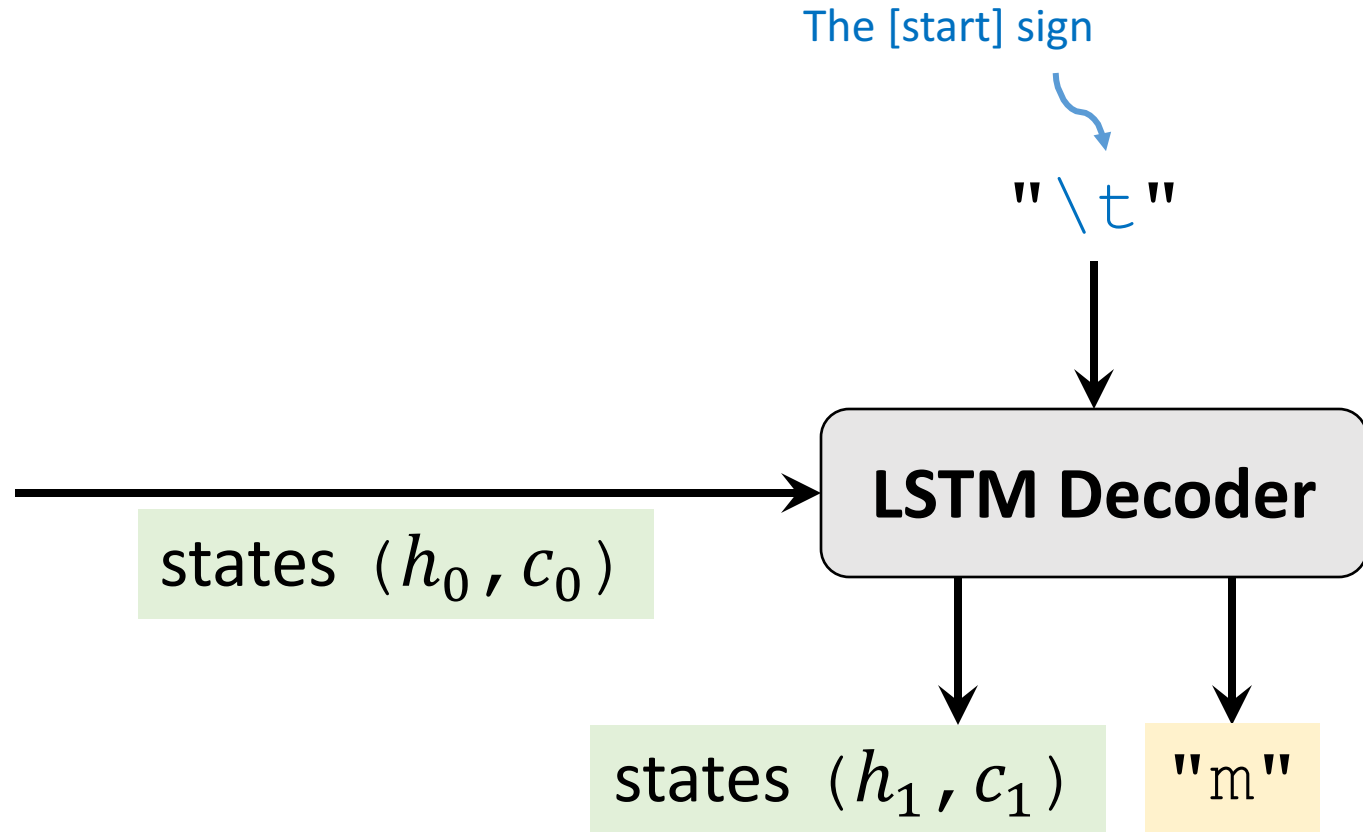
Inference



Inference

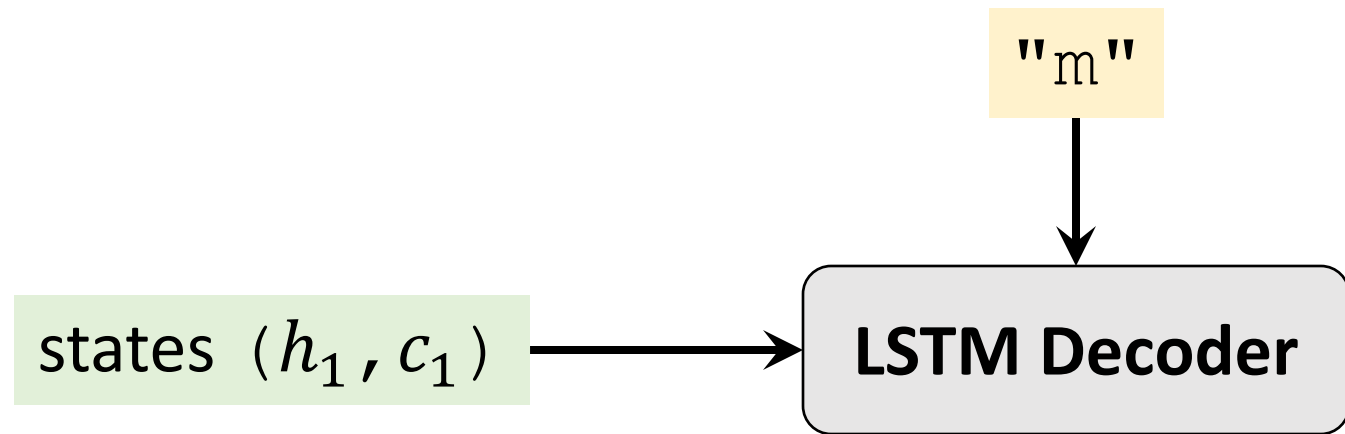


Inference



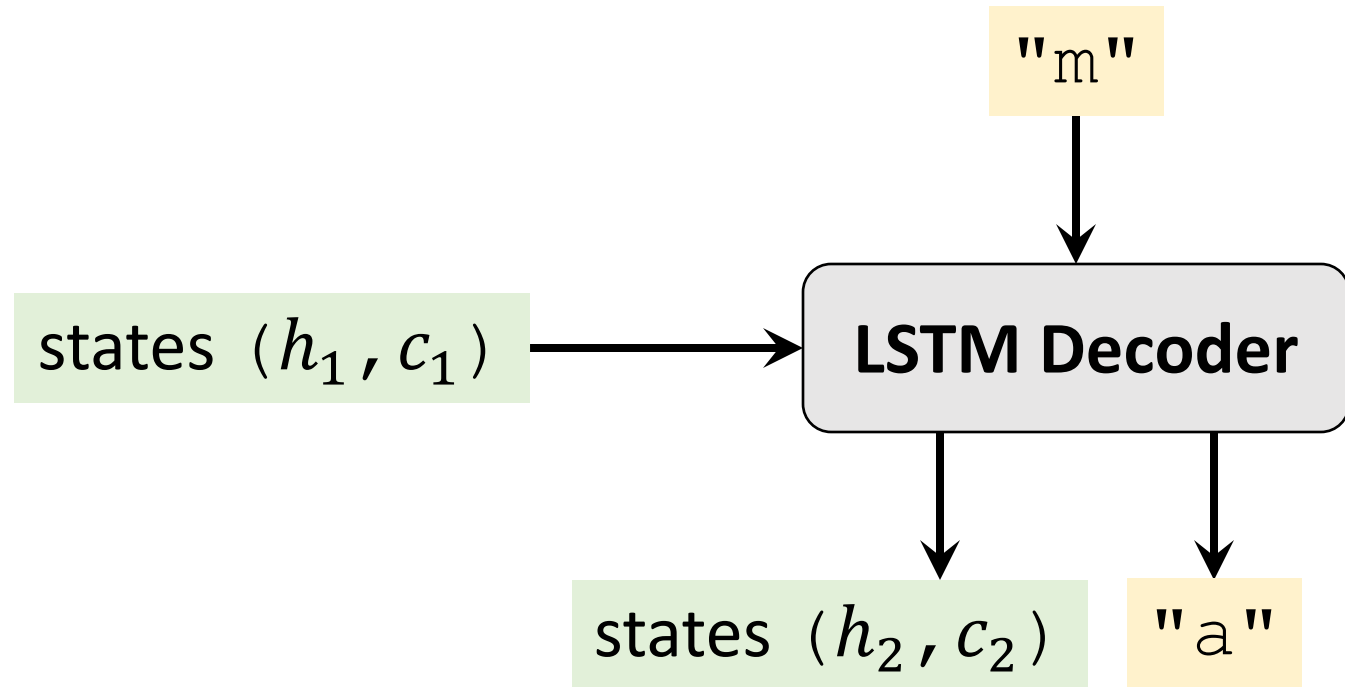
Record: `m`

Inference



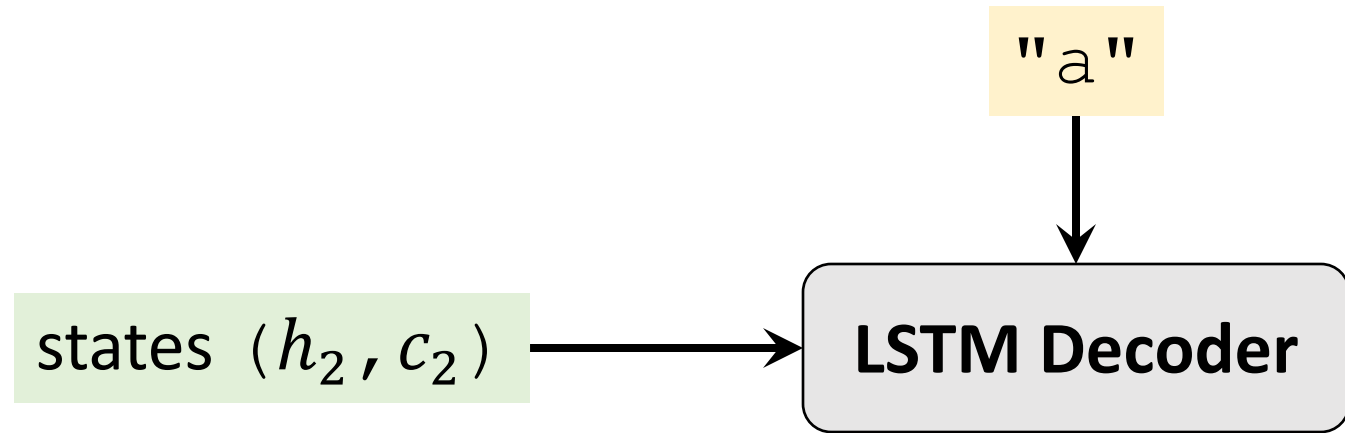
Record: " m "

Inference



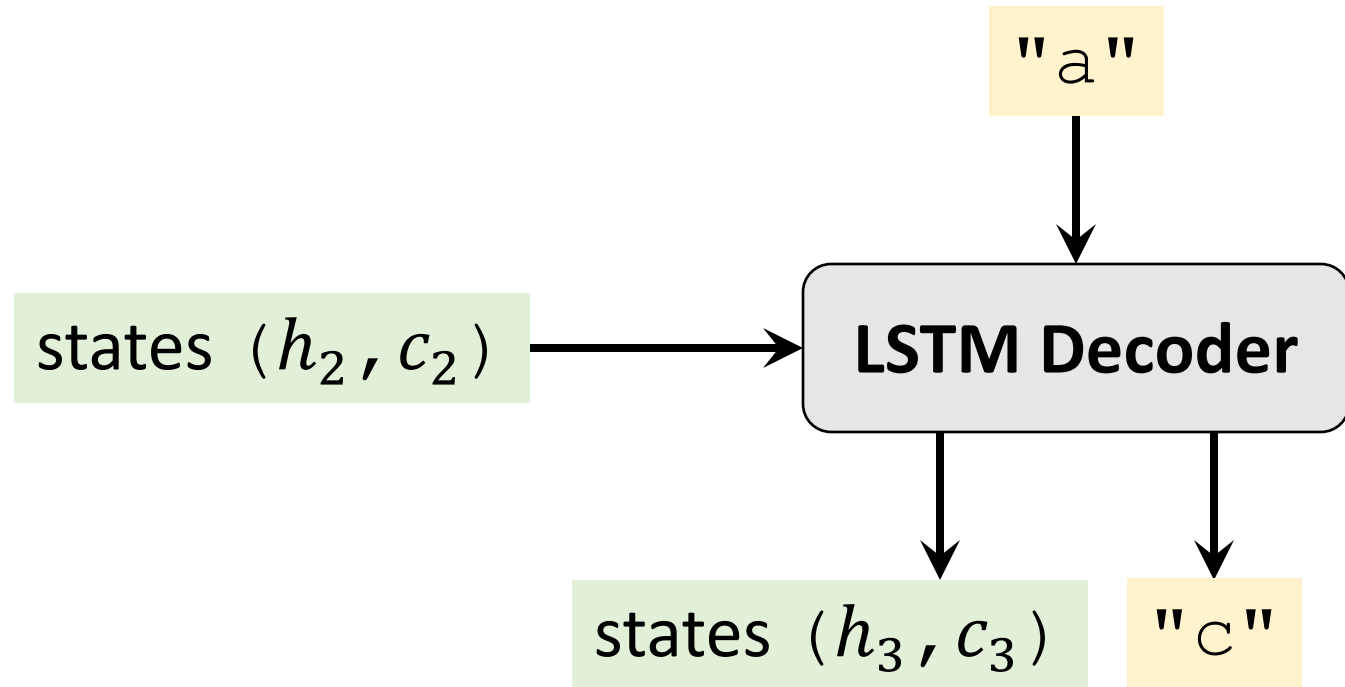
Record: "ma"

Inference



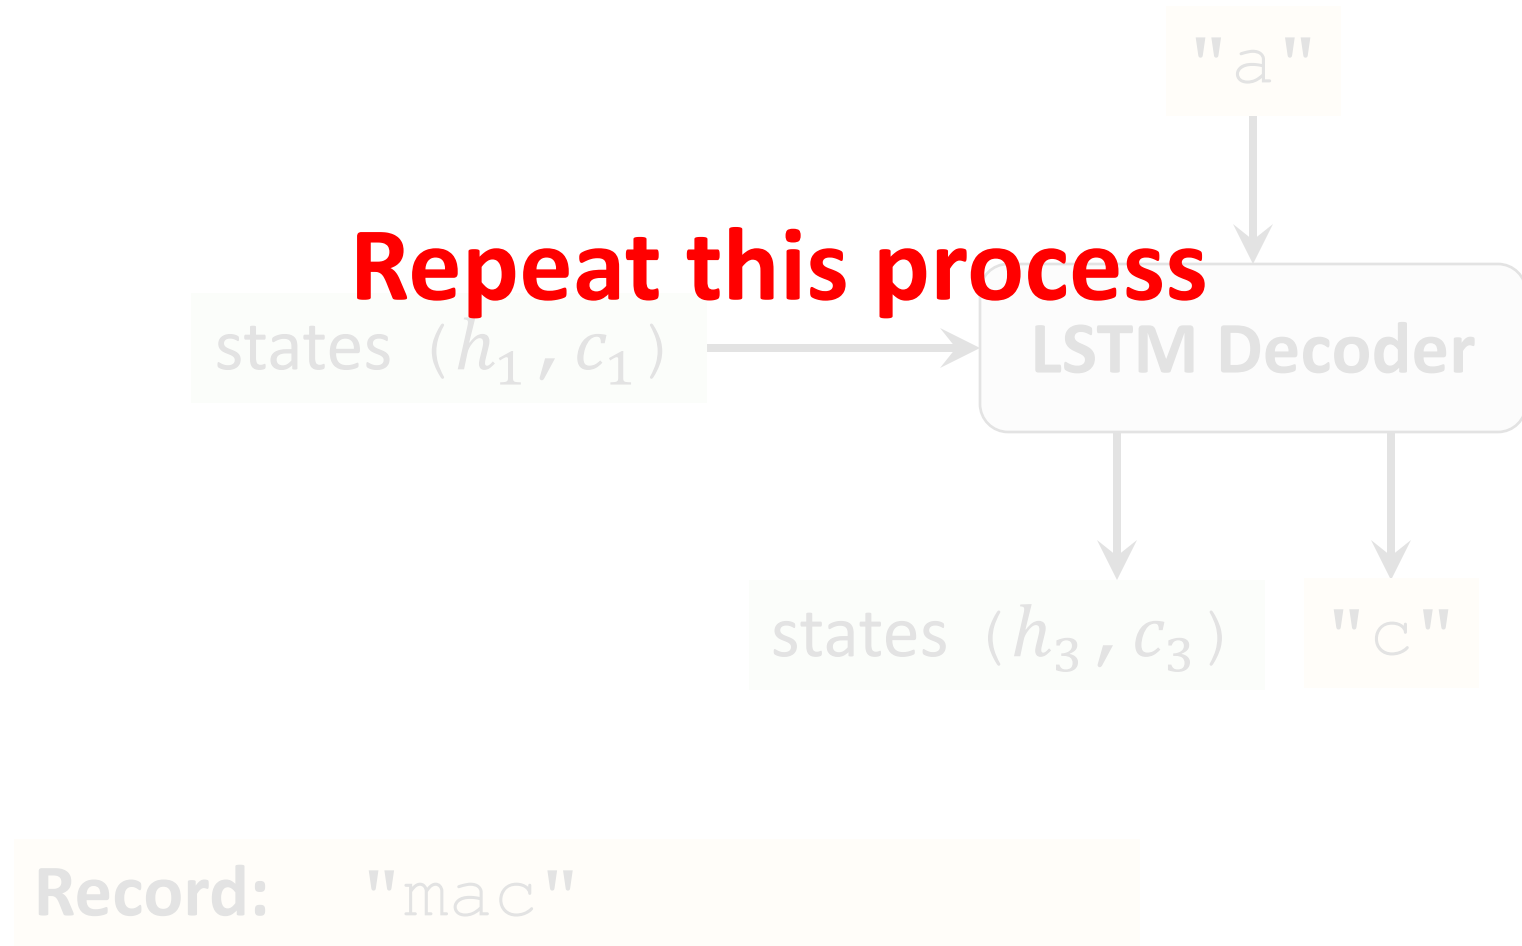
Record: "ma"

Inference

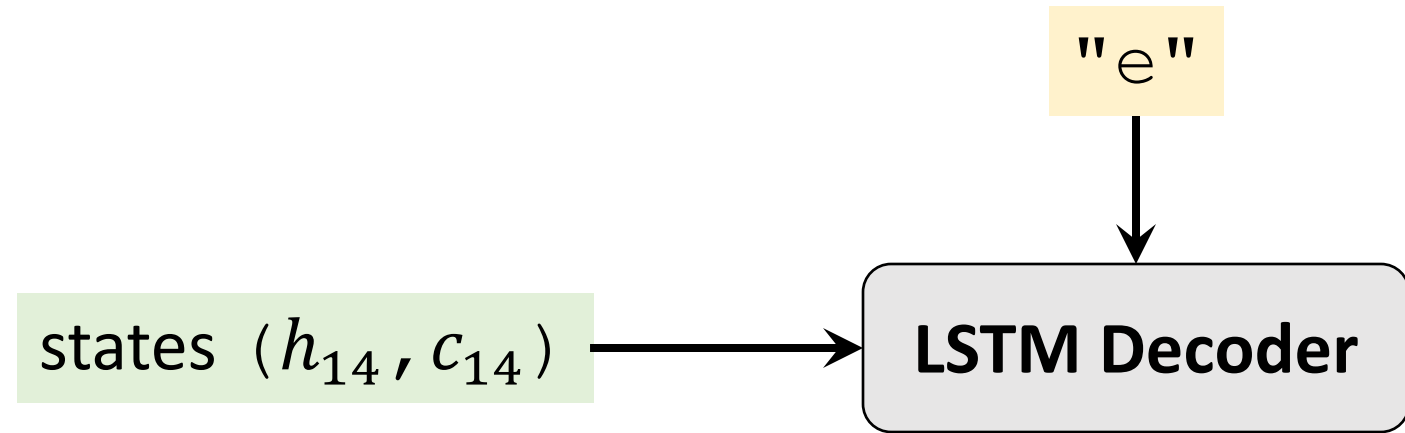


Record: "mac"

Inference

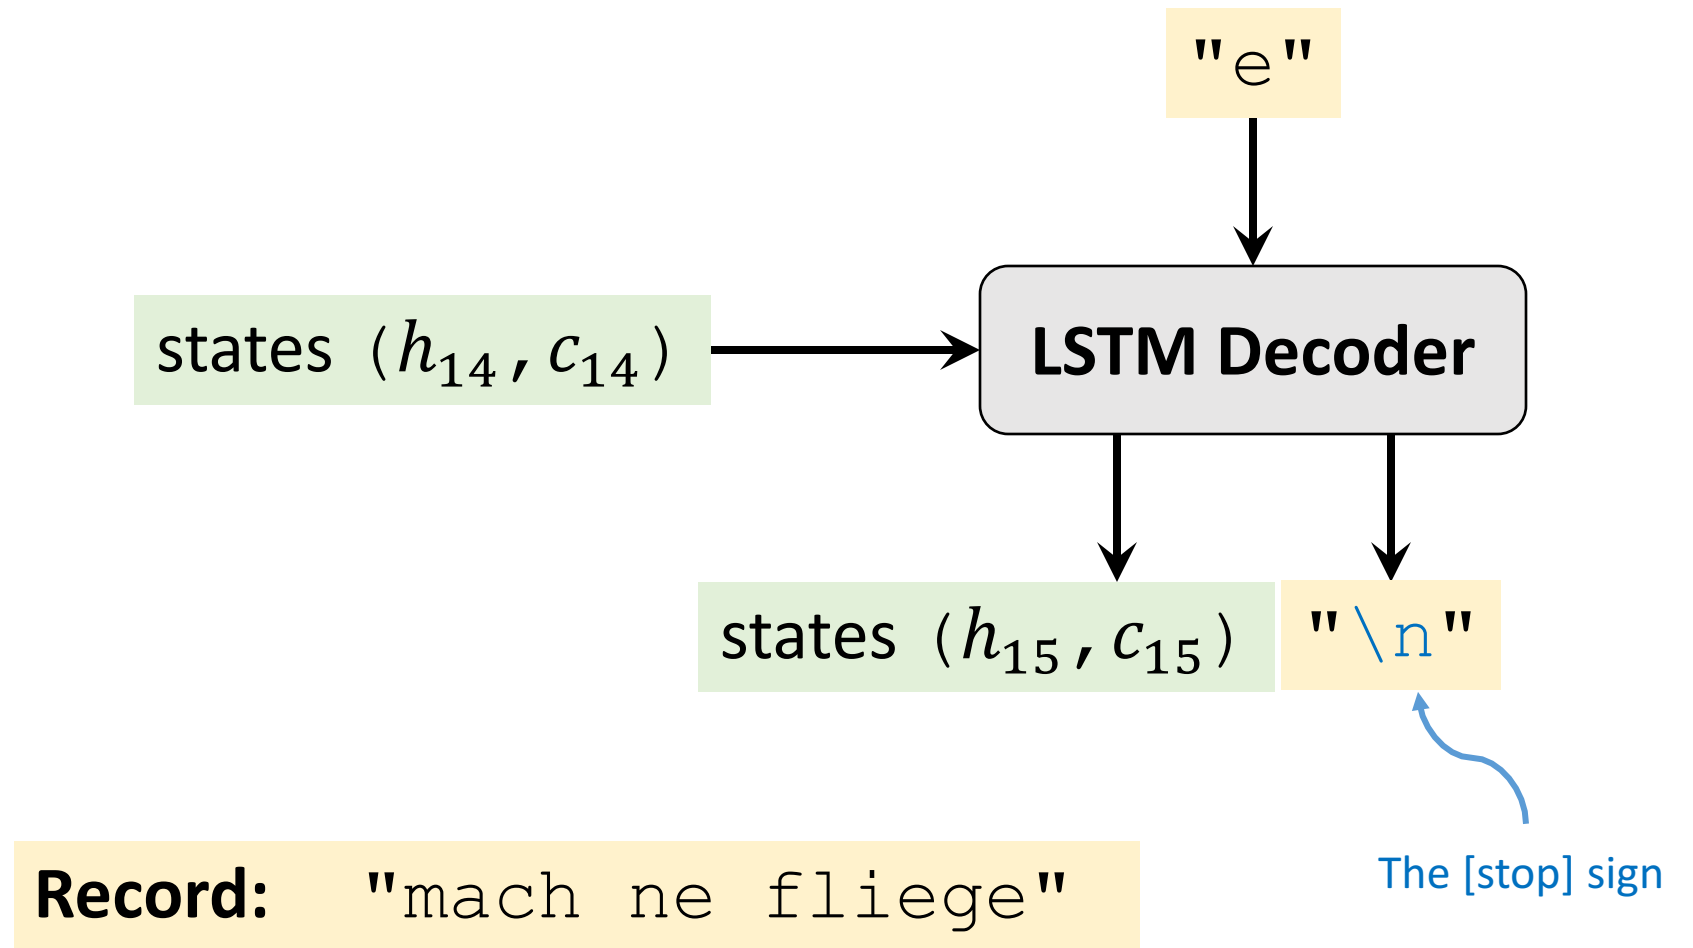


Inference

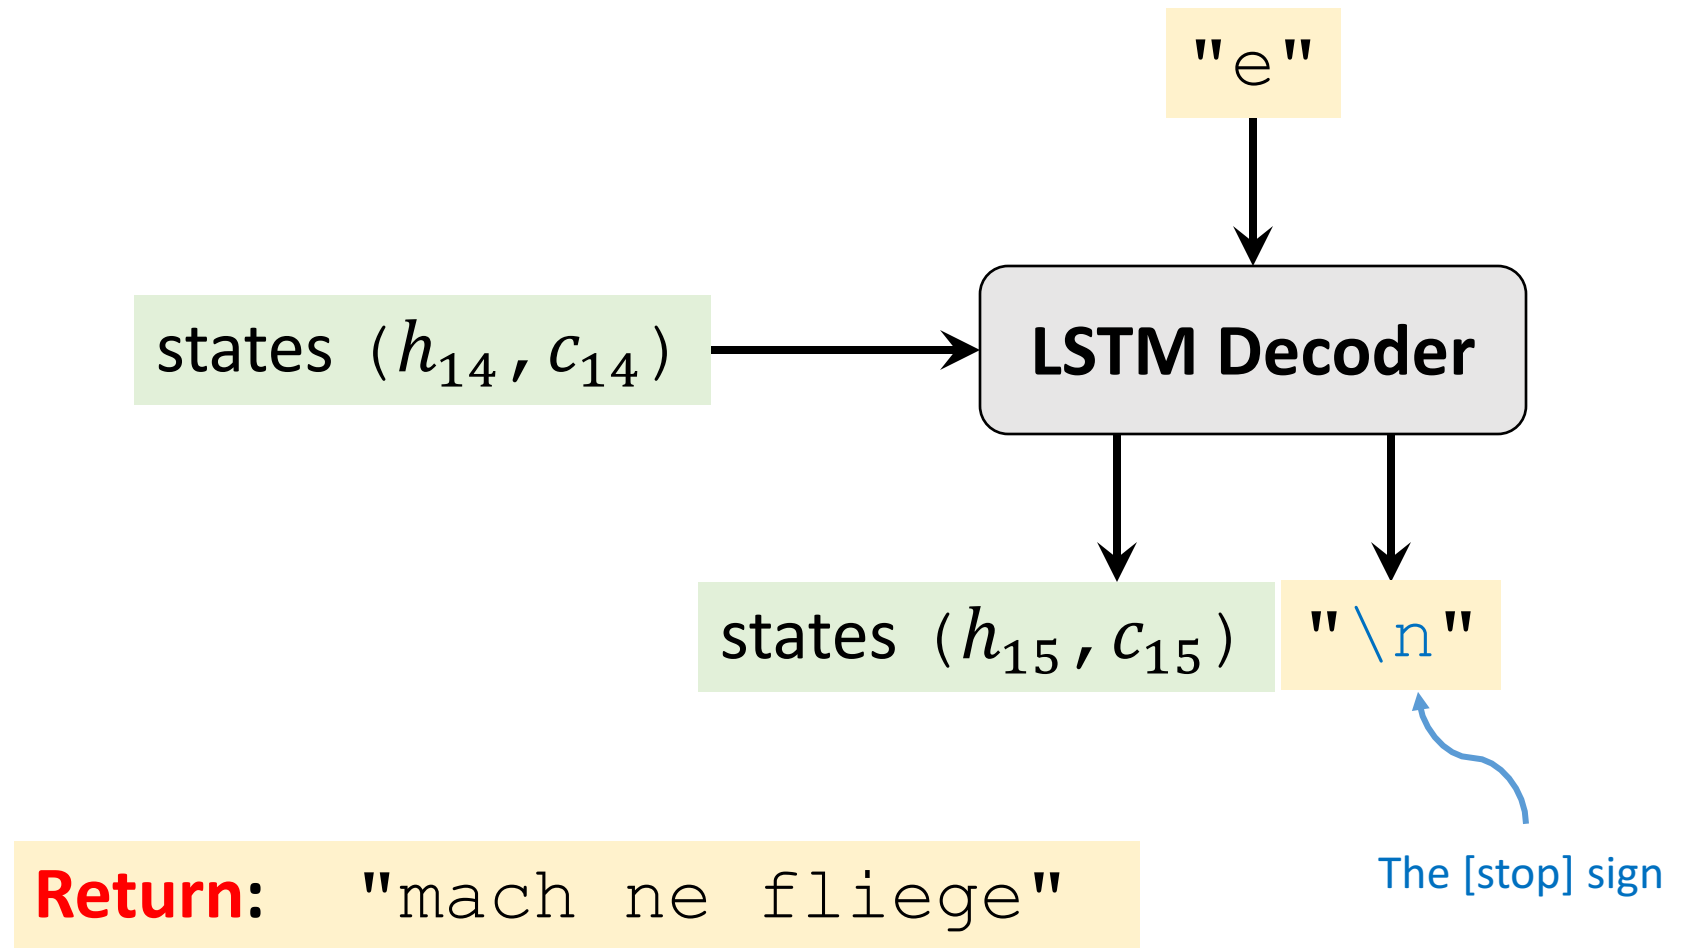


Record: "mach ne fliege"

Inference



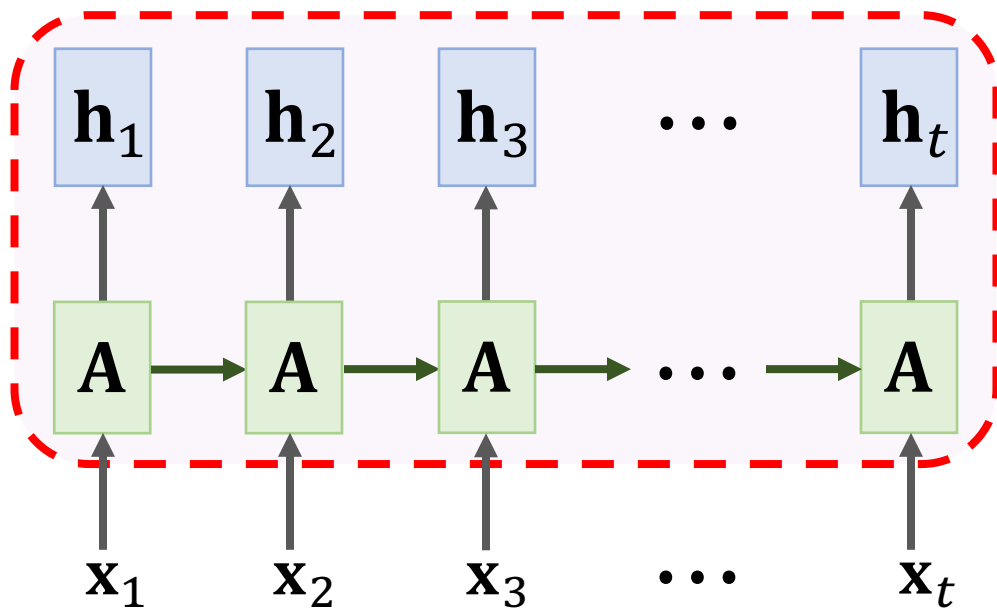
Inference



Summary

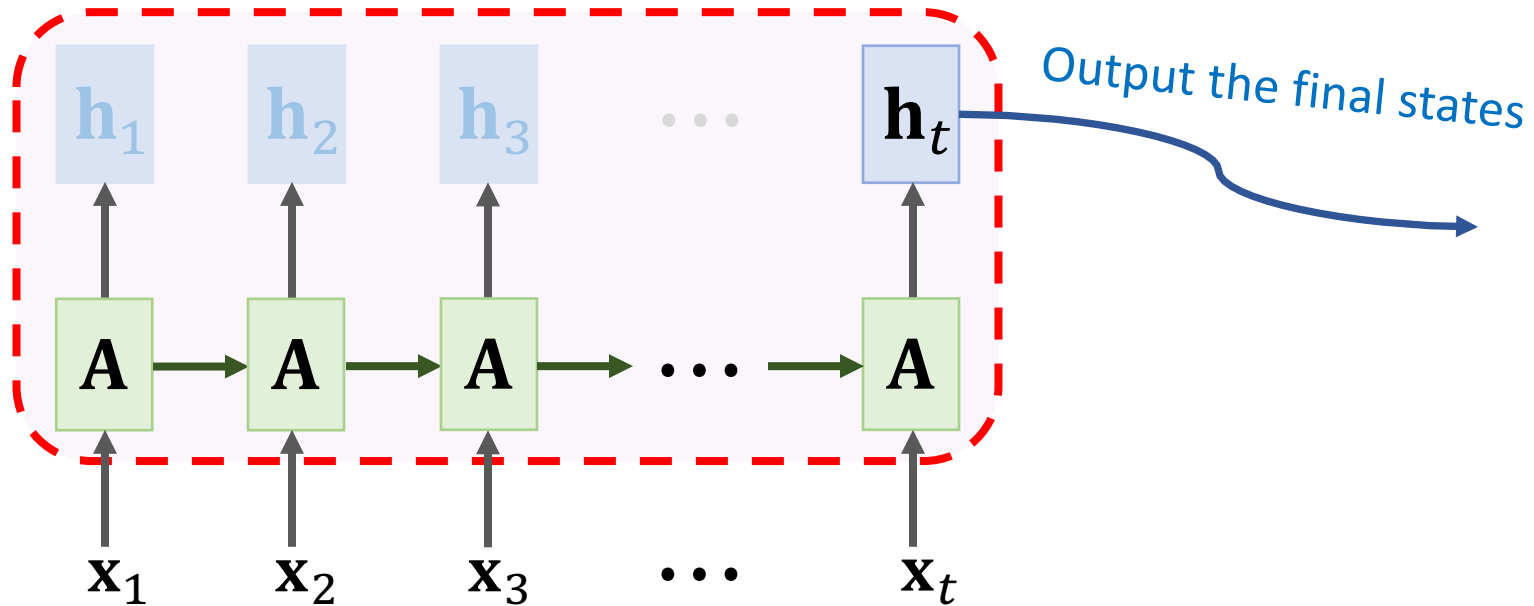
Seq2Seq Model

Encoder RNN



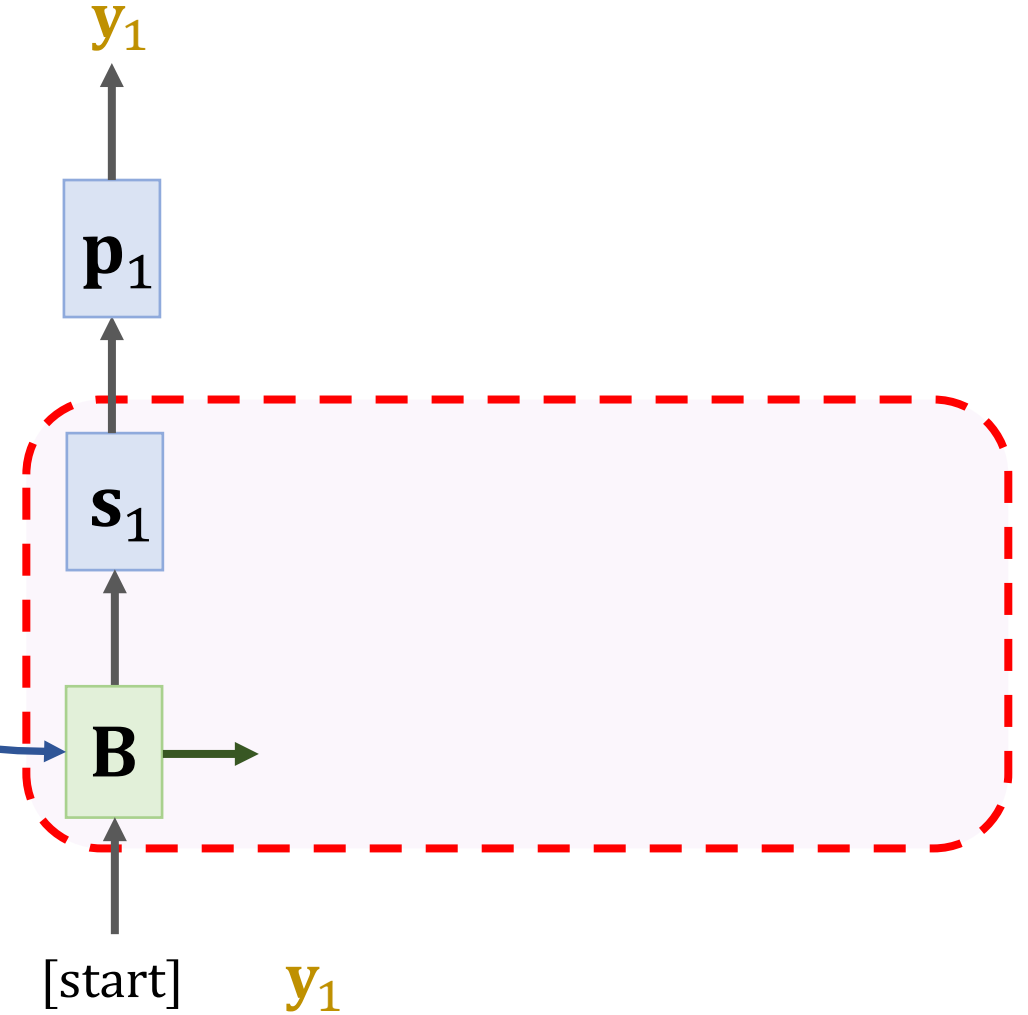
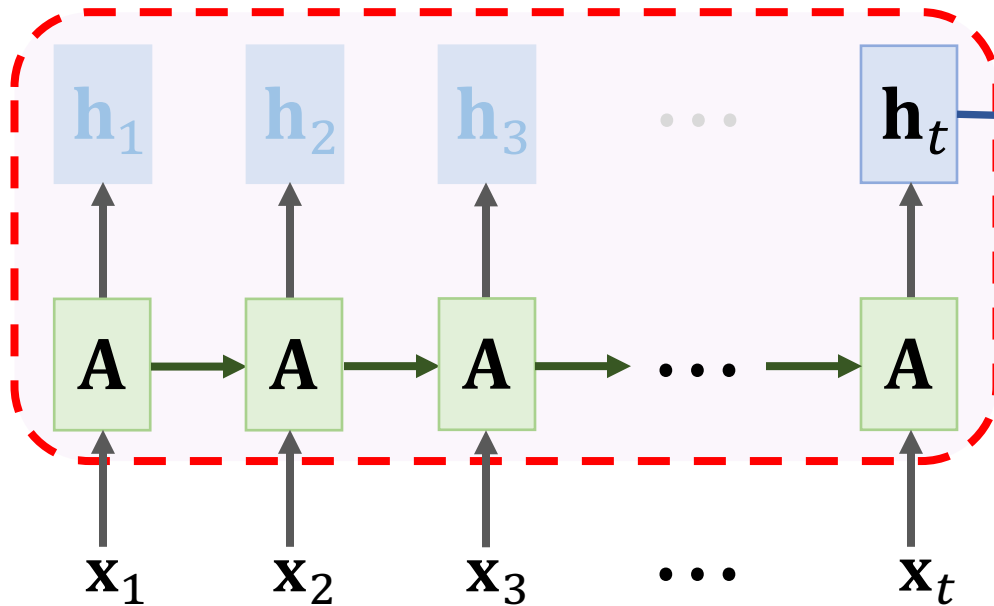
Seq2Seq Model

Encoder RNN



Seq2Seq Model

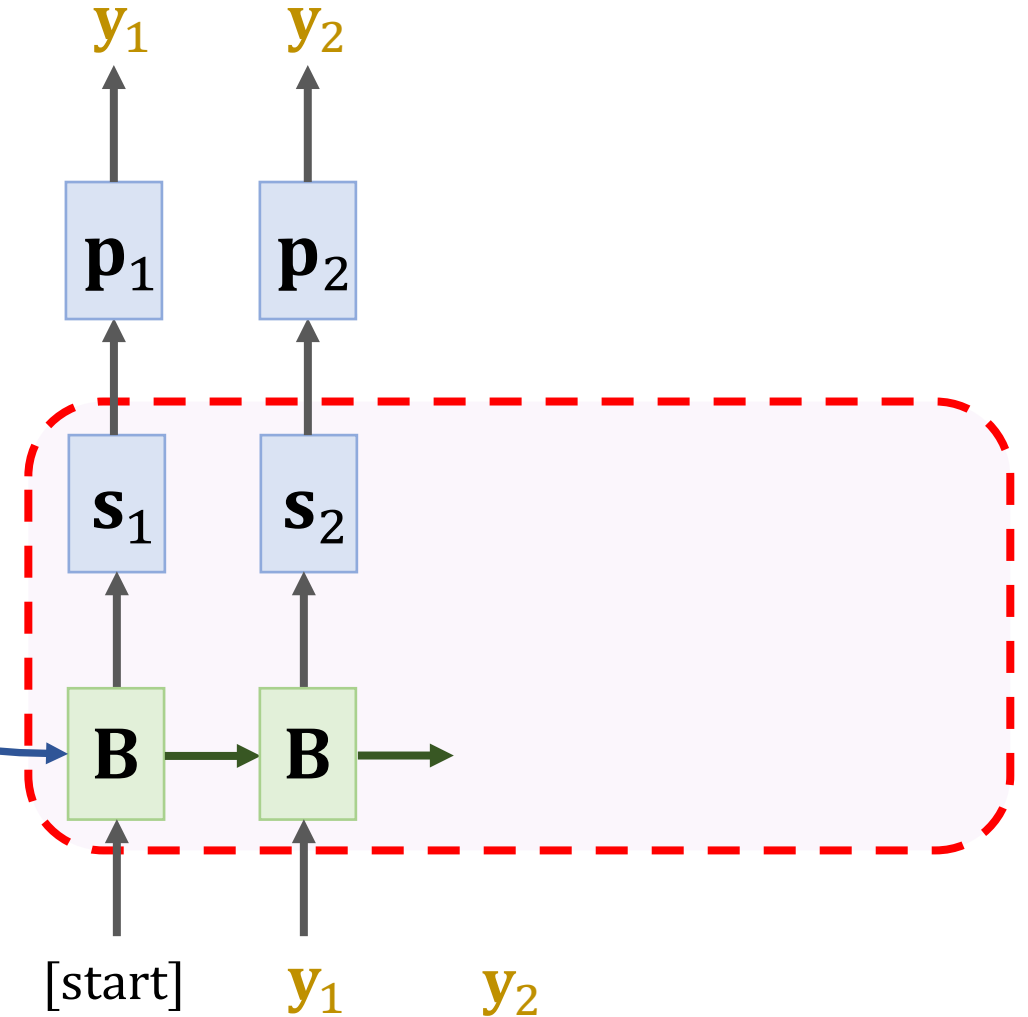
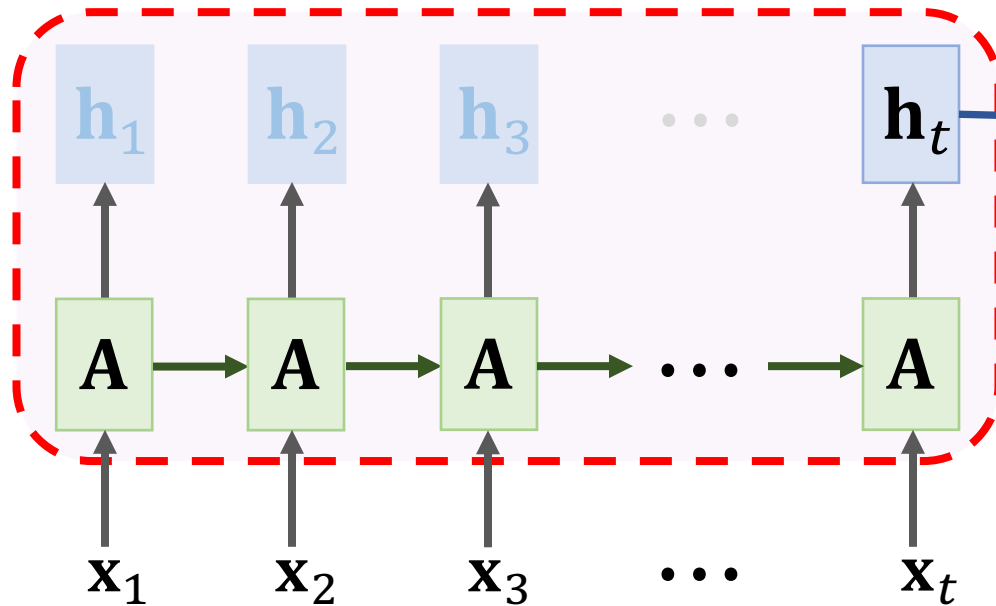
Encoder RNN



Decoder RNN

Seq2Seq Model

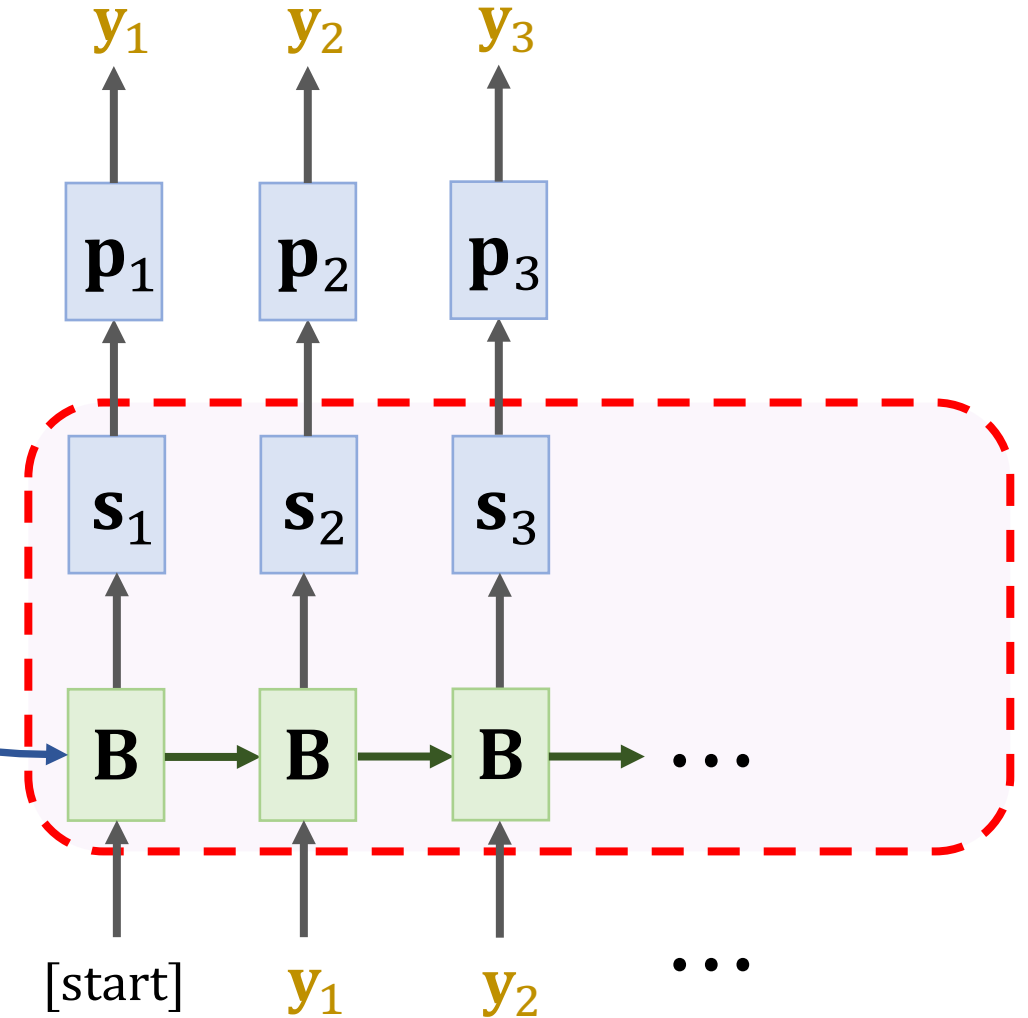
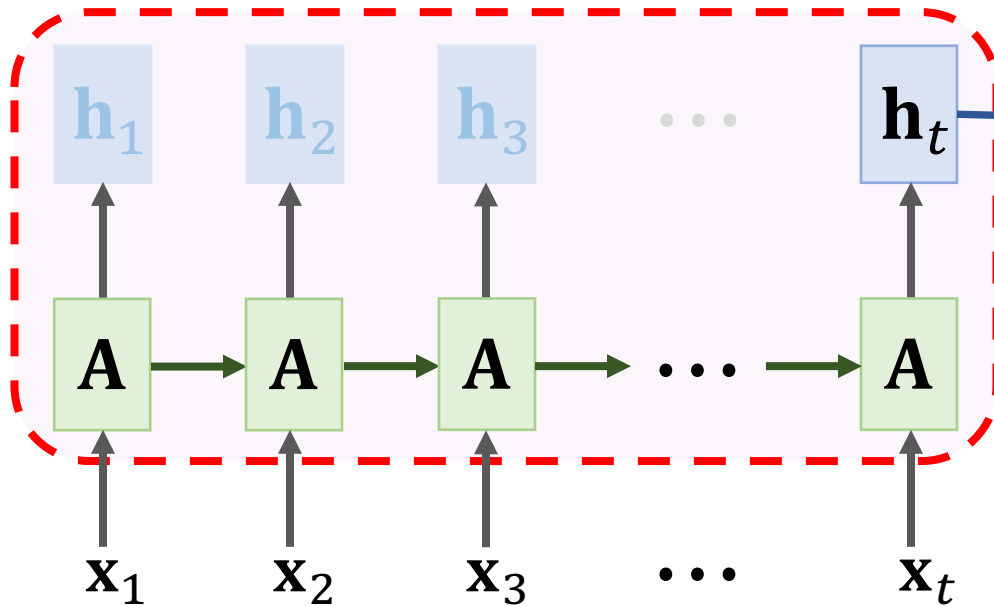
Encoder RNN



Decoder RNN

Seq2Seq Model

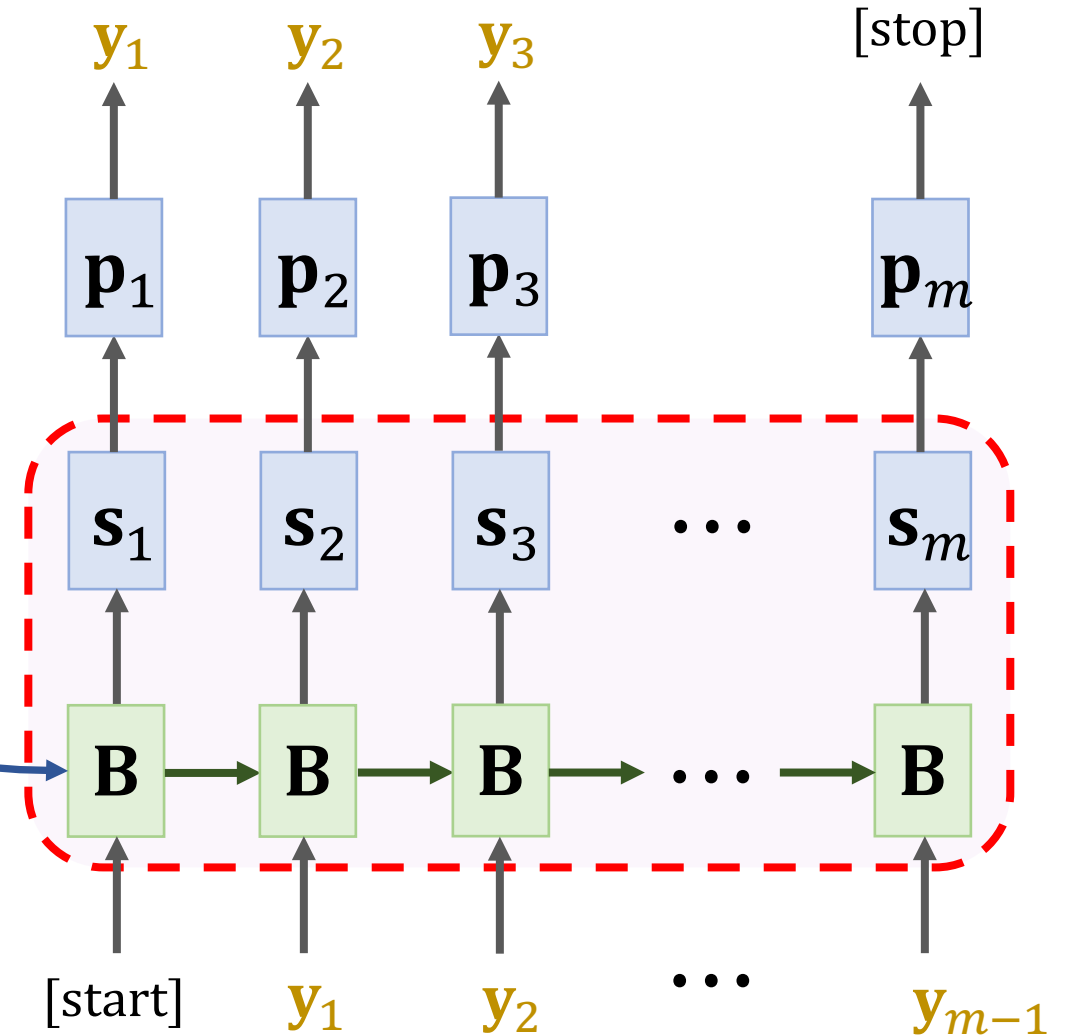
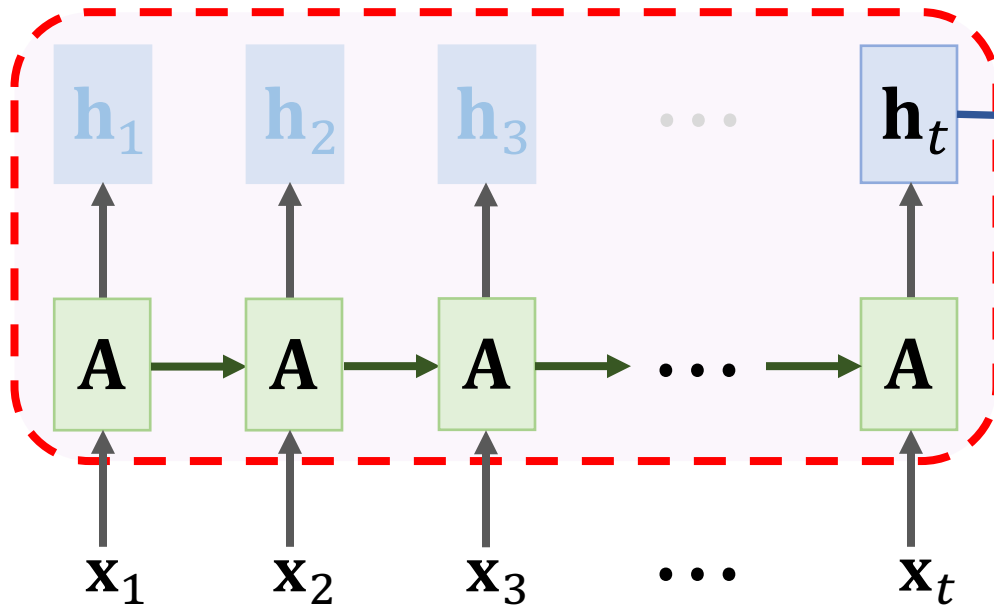
Encoder RNN



Decoder RNN

Seq2Seq Model

Encoder RNN

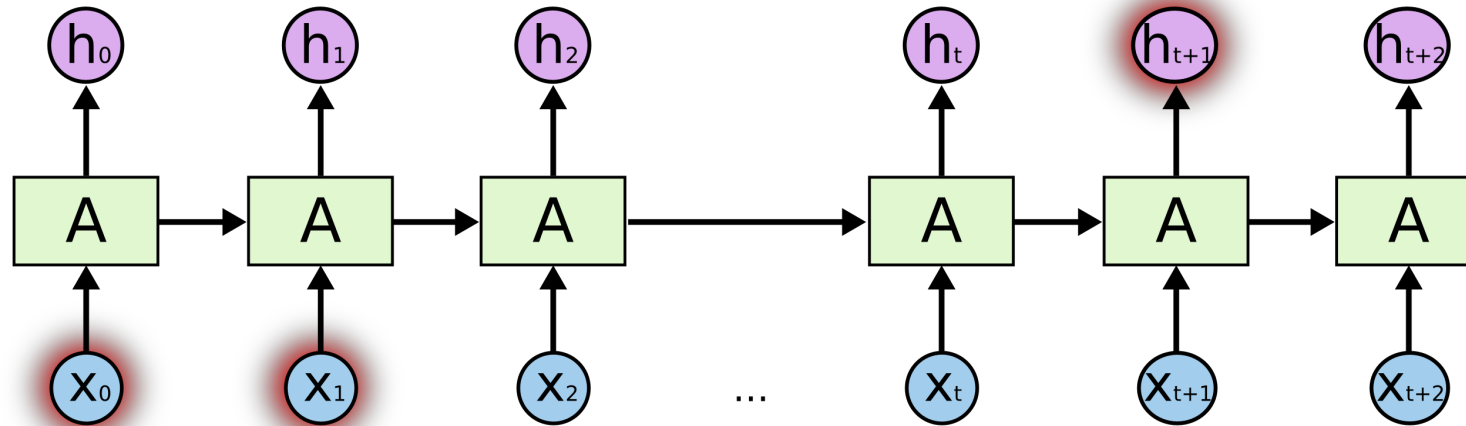


Decoder RNN

How to Improve?

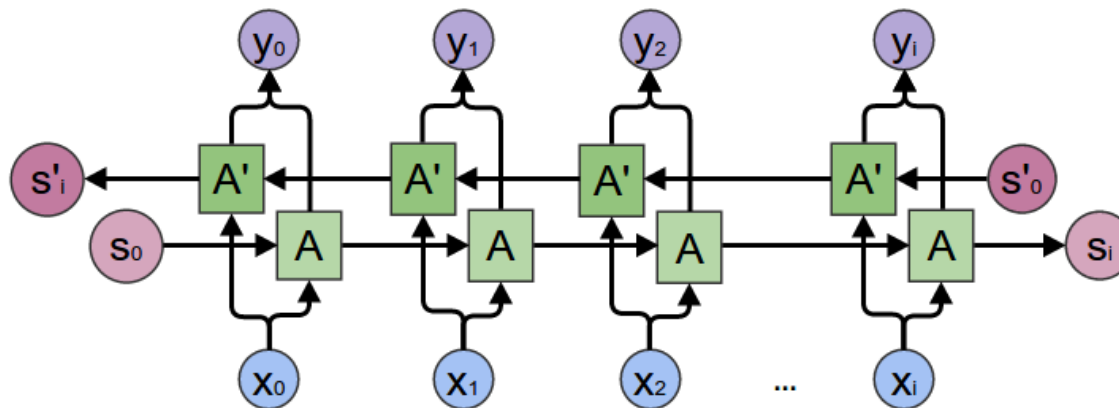
1. Bi-LSTM instead of LSTM

- The final states (\mathbf{h}_t and \mathbf{c}_t) of the Encoder have all the information of the English sentence.
- Really?
 - If the sentence is long (hundreds of tokens), the final states have forgotten the first tokens.



1. Bi-LSTM instead of LSTM

- The final states (\mathbf{h}_t and \mathbf{c}_t) of the Encoder have all the information of the English sentence.
- Really?
 - If the sentence is long (hundreds of tokens), the final states have forgotten the first tokens.
- Bi-LSTM (left-to-right and right-to-left) remembers the first tokens.



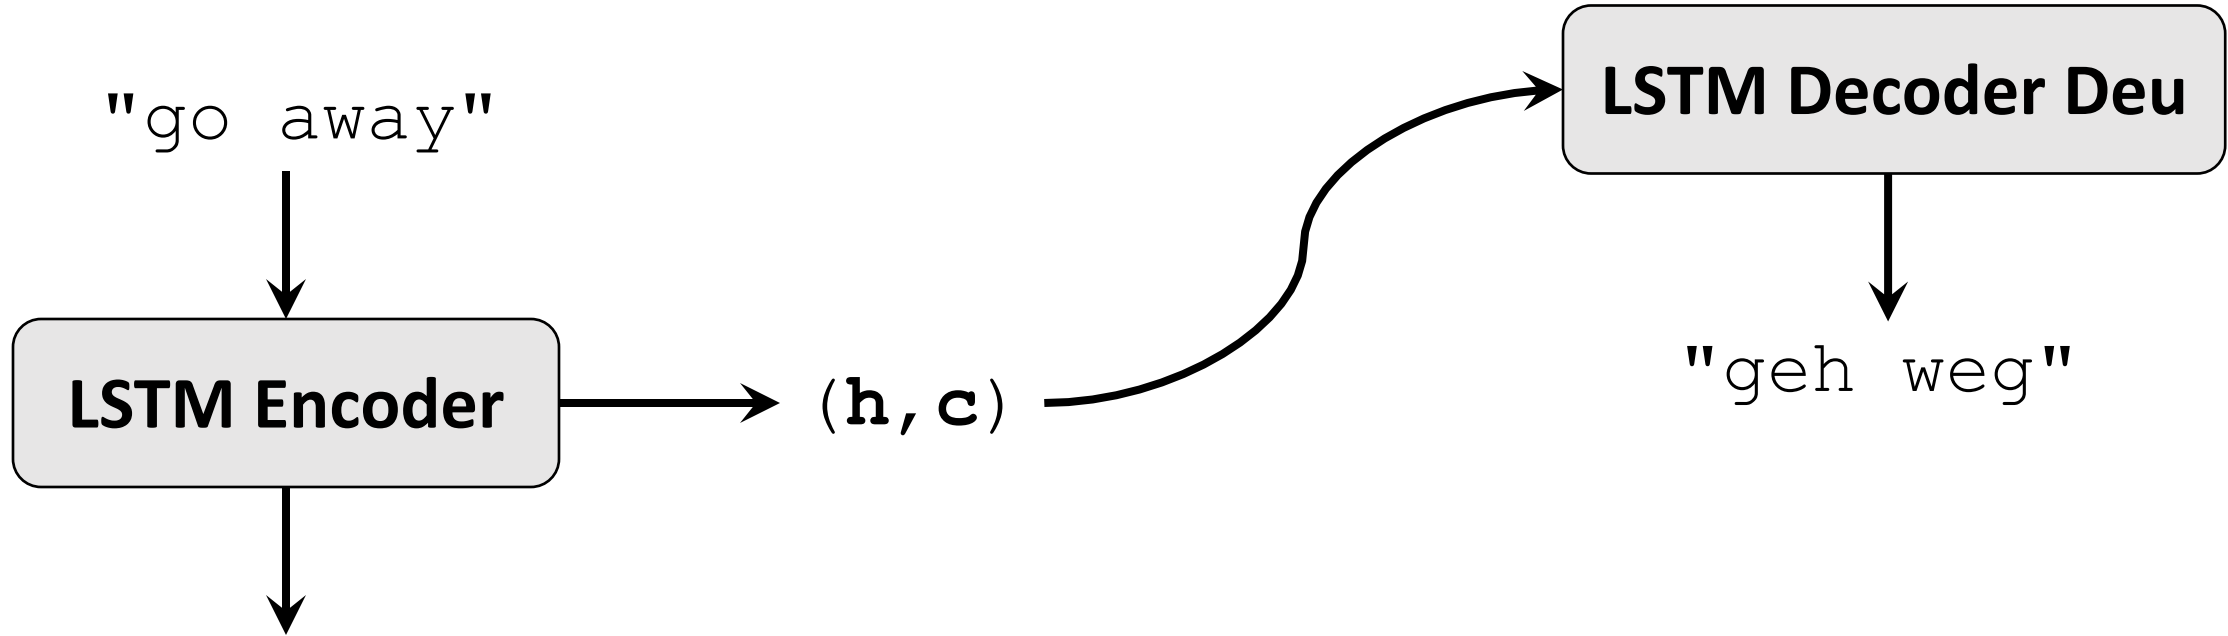
2. Word-Level Tokenization

- Word-level tokenization instead of char-level.
 - The average length of English words is 4.5 letters.
 - The sequences will be 4.5x shorter.
 - Shorter sequence, less likely to forget.

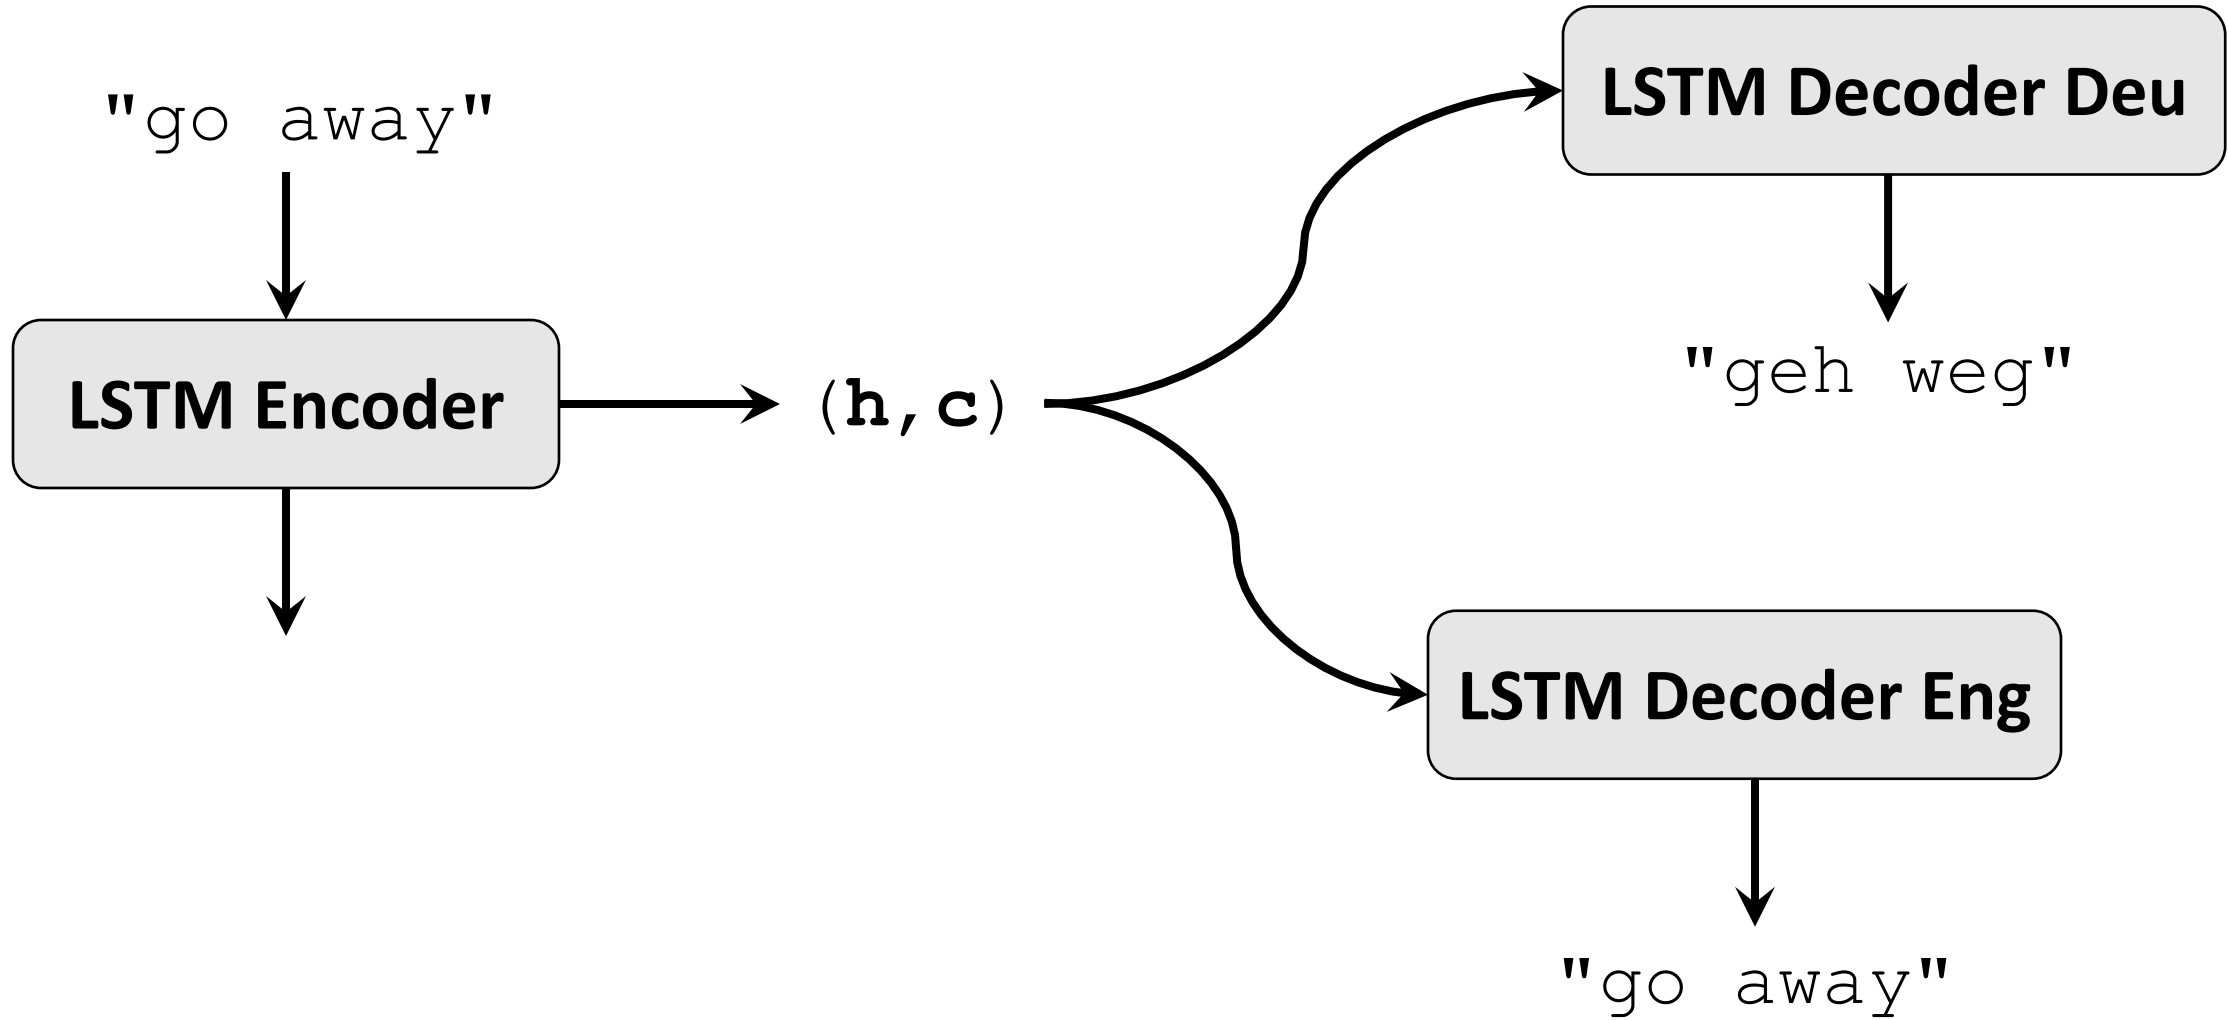
2. Word-Level Tokenization

- Word-level tokenization instead of char-level.
 - The average length of English words is 4.5 letters.
 - The sequences will be 4.5x shorter.
 - Shorter sequence, less likely to forget.
- But you will need a large dataset!
 - # of (frequently used) chars is $\sim 10^2$.
 - # of (frequently used) words is $\sim 10^4$.
 - Word Embedding Layer has many parameters → overfitting!

3. Multi-Task Learning



3. Multi-Task Learning



3. Multi-Task Learning

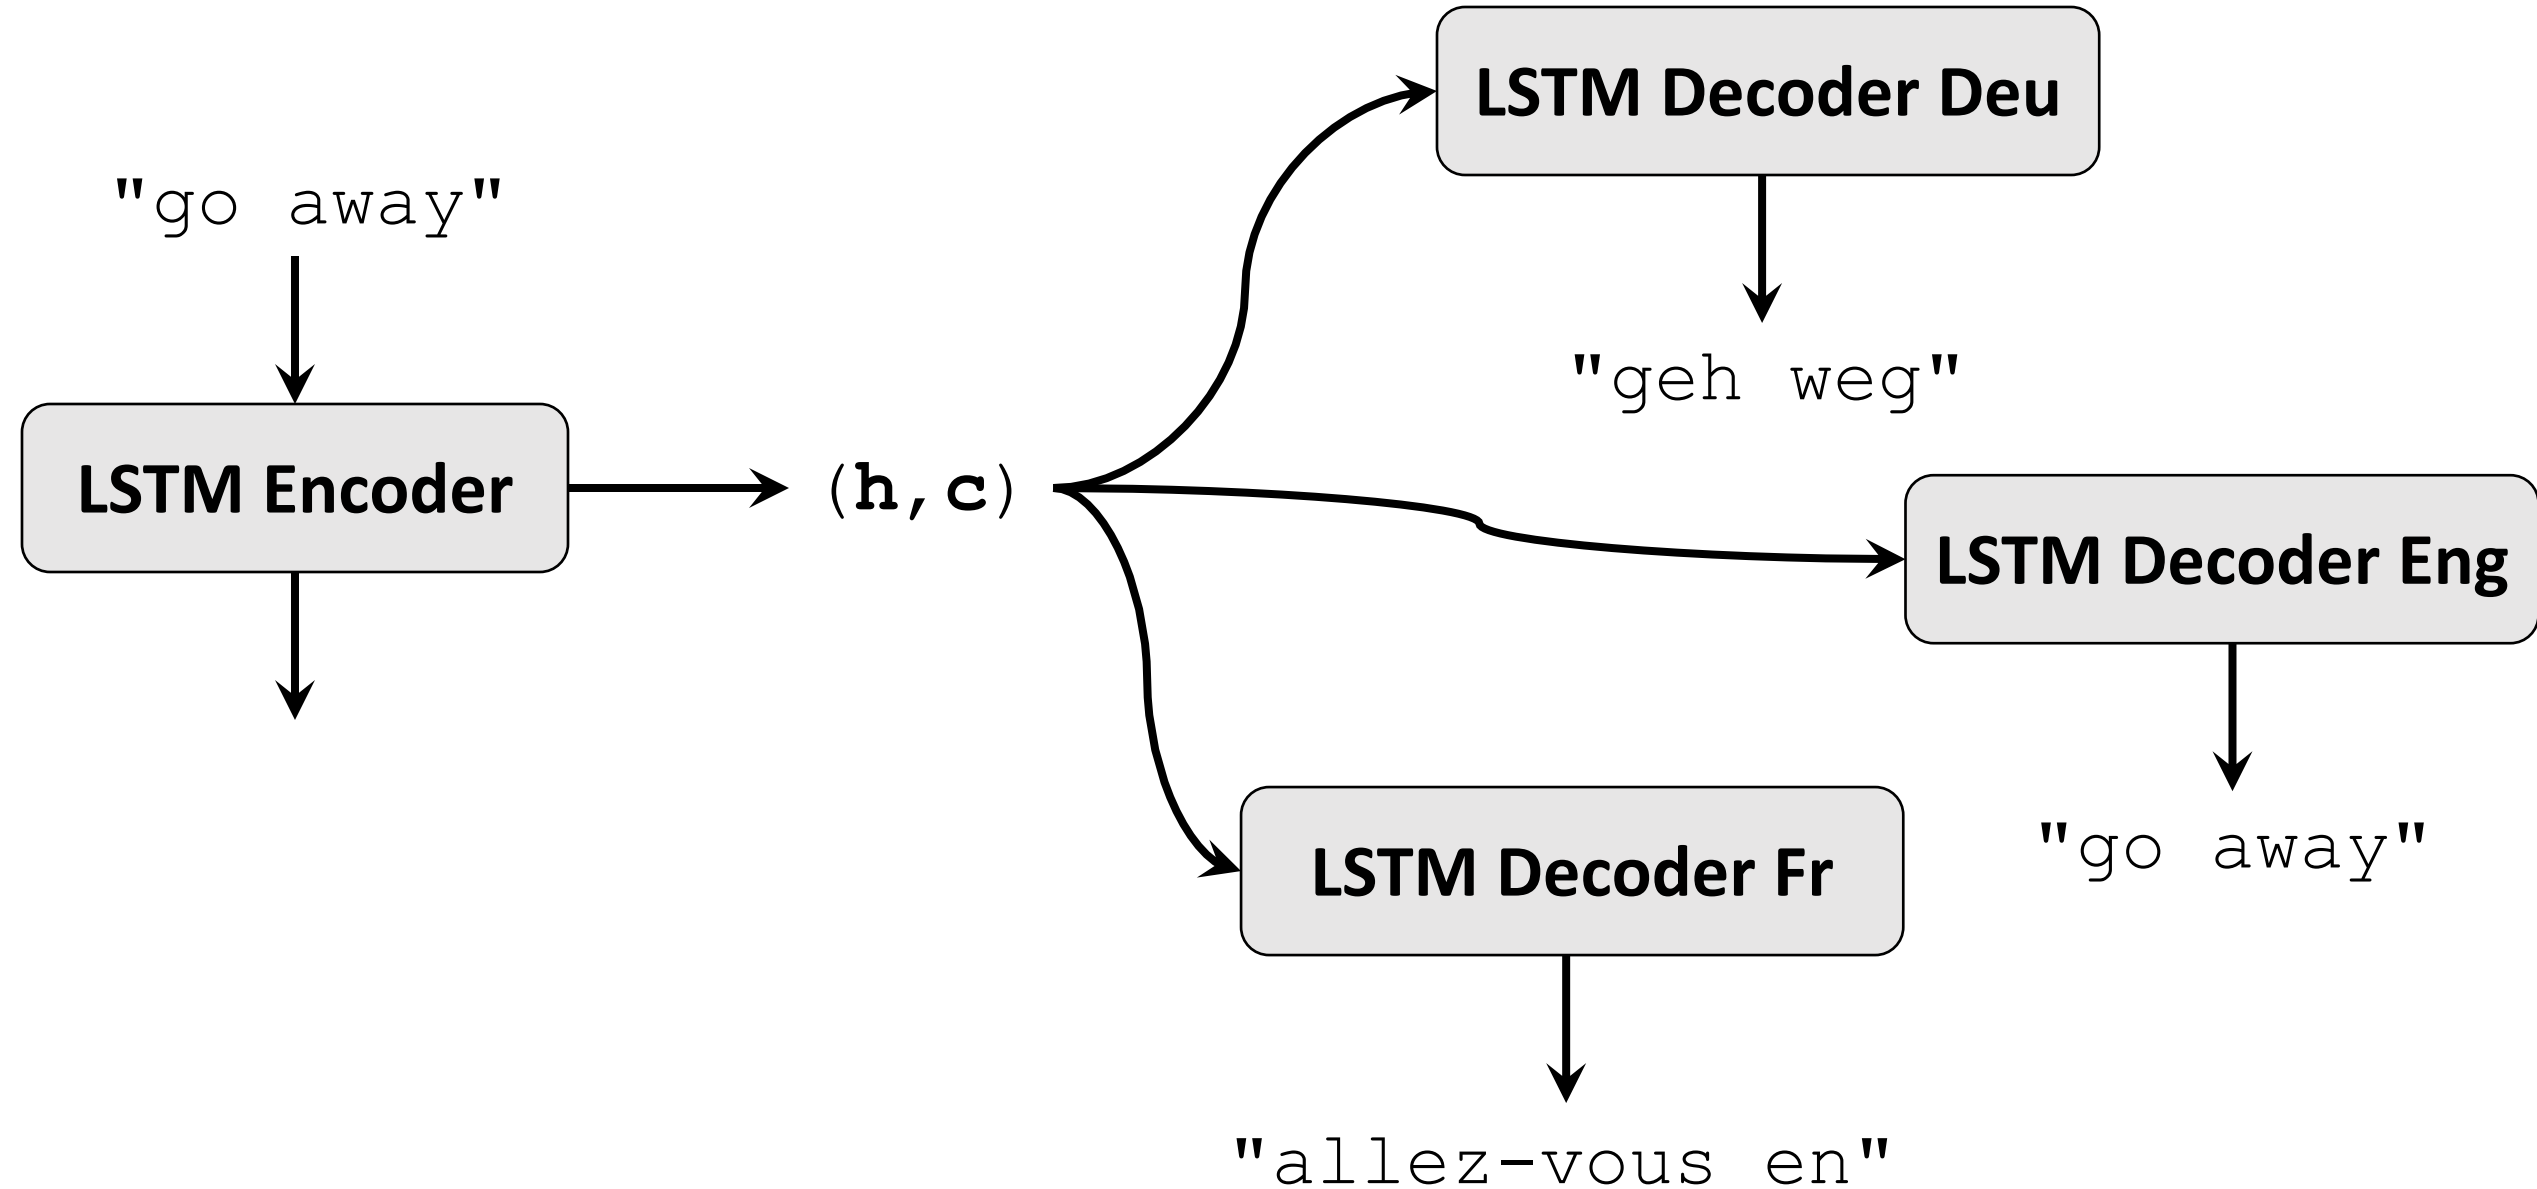
- Even if you want to **translate English to German**, you can use all the datasets:

-  Afrikaans - English [afr-eng.zip](#) (725)
-  Aklanon - English [akl-eng.zip](#) (22)
-  Albanian - English [sqi-eng.zip](#) (412)
-  Algerian Arabic - English [arq-eng.zip](#) (156)
-  Arabic - English [ara-eng.zip](#) (11009)
-  Arabic (Gulf) - English [afb-eng.zip](#) (28)
-  Assamese - English [asm-eng.zip](#) (23)
-  Asturian - English [ast-eng.zip](#) (23)
-  Azerbaijani - English [aze-eng.zip](#) (2131)
-  Basque - English [eus-eng.zip](#) (667)
-  Belarusian - English [bel-eng.zip](#) (2698)
-  Bengali - English [ben-eng.zip](#) (4399)
-  Berber - English [ber-eng.zip](#) (54988)
-  Bulgarian - English [bul-eng.zip](#) (14968)

• • •

-  Spanish - English [spa-eng.zip](#) (120799)
-  Swedish - English [swe-eng.zip](#) (17409)
-  Tagalog - English [tgl-eng.zip](#) (3144)
-  Tamil - English [tam-eng.zip](#) (197)
-  Tatar - English [tat-eng.zip](#) (529)
-  Telugu - English [tel-eng.zip](#) (138)
-  Thai - English [tha-eng.zip](#) (110)
-  Turkish - English [tur-eng.zip](#) (497249)
-  Ukrainian - English [ukr-eng.zip](#) (113396)
-  Urdu - English [urd-eng.zip](#) (1180)
-  Uyghur - English [uig-eng.zip](#) (285)
-  Vietnamese - English [vie-eng.zip](#) (3413)
-  Waray - English [war-eng.zip](#) (1181)
-  Zaza - English [zza-eng.zip](#) (345)

3. Multi-Task Learning



How to Improve?

1. Bi-LSTM instead of LSTM.
2. Tokenization in the word-level (rather than char-level.)
3. Multi-task learning.
4. Attention! (Future lectures, soon.)

Homework 4

- Build a seq2seq model for machine translation.
 - Anything languages except for [English ==> German].
 - Follow my IPython Notebook.
- Make as least one improvement over my naïve model.
 - E.g., Bi-LSTM, multi-task, attention, etc.
- Evaluate your model using BLEU score. (Optional.)
 - BLEU (BiLingual Evaluation Understudy).
 - Reference:
 - Wikipedia: <https://en.wikipedia.org/wiki/BLEU>
 - Blog: <https://machinelearningmastery.com/calculate-bleu-score-for-text-python/>